

Lecture 16

A Brief Introduction to Generative Adversarial Networks (GAN)

STAT 453: Deep Learning, Spring 2020

Sebastian Raschka

pages.stat.wisc.edu/~sraschka/teaching/stat453-ss2020/

Lecture material:

<https://github.com/rasbt/stat453-deep-learning-ss20/tree/master/L16-gan>

Zhu, J. Y., Park, T., Isola, P., & Efros, A. A. (2017). Unpaired image-to-image translation using cycle-consistent adversarial networks. In *Proceedings of the IEEE International Conference on Computer Vision* (pp. 2223-2232).

<https://arxiv.org/abs/1703.10593>



<https://github.com/junyanz/CycleGAN>

<https://www.youtube.com/watch?v=9reHvktowLY>

Zhu, J. Y., Park, T., Isola, P., & Efros, A. A. (2017). Unpaired image-to-image translation using cycle-consistent adversarial networks. In *Proceedings of the IEEE International Conference on Computer Vision* (pp. 2223-2232).

<https://arxiv.org/abs/1703.10593>



<https://github.com/junyanz/CycleGAN>

<https://www.youtube.com/watch?v=9reHvktoV>

Question: Why does the model change the background as well?



<https://thispersondoesnotexist.com>

Karras, T., Laine, S., & Aila, T. (2019). A style-based generator architecture for generative adversarial networks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition* (pp. 4401-4410).

<https://arxiv.org/abs/1812.04948>

Click on the person who is real.



<http://www.whichfaceisreal.com>

Lecture Overview

- 1. The Main Idea Behind GAN**
2. The GAN Training
3. GANs in Practice
4. GAN Resources

Generative Adversarial Networks (GAN)

- The original purpose is to generate new data
- Classically for generating new images, but applicable to wide range of domains
- Learns the training set distribution and can generate new images that have never been seen before
- In contrast to e.g., autoregressive models or RNNs (generating one word at a time), GANs generate the whole output all at once

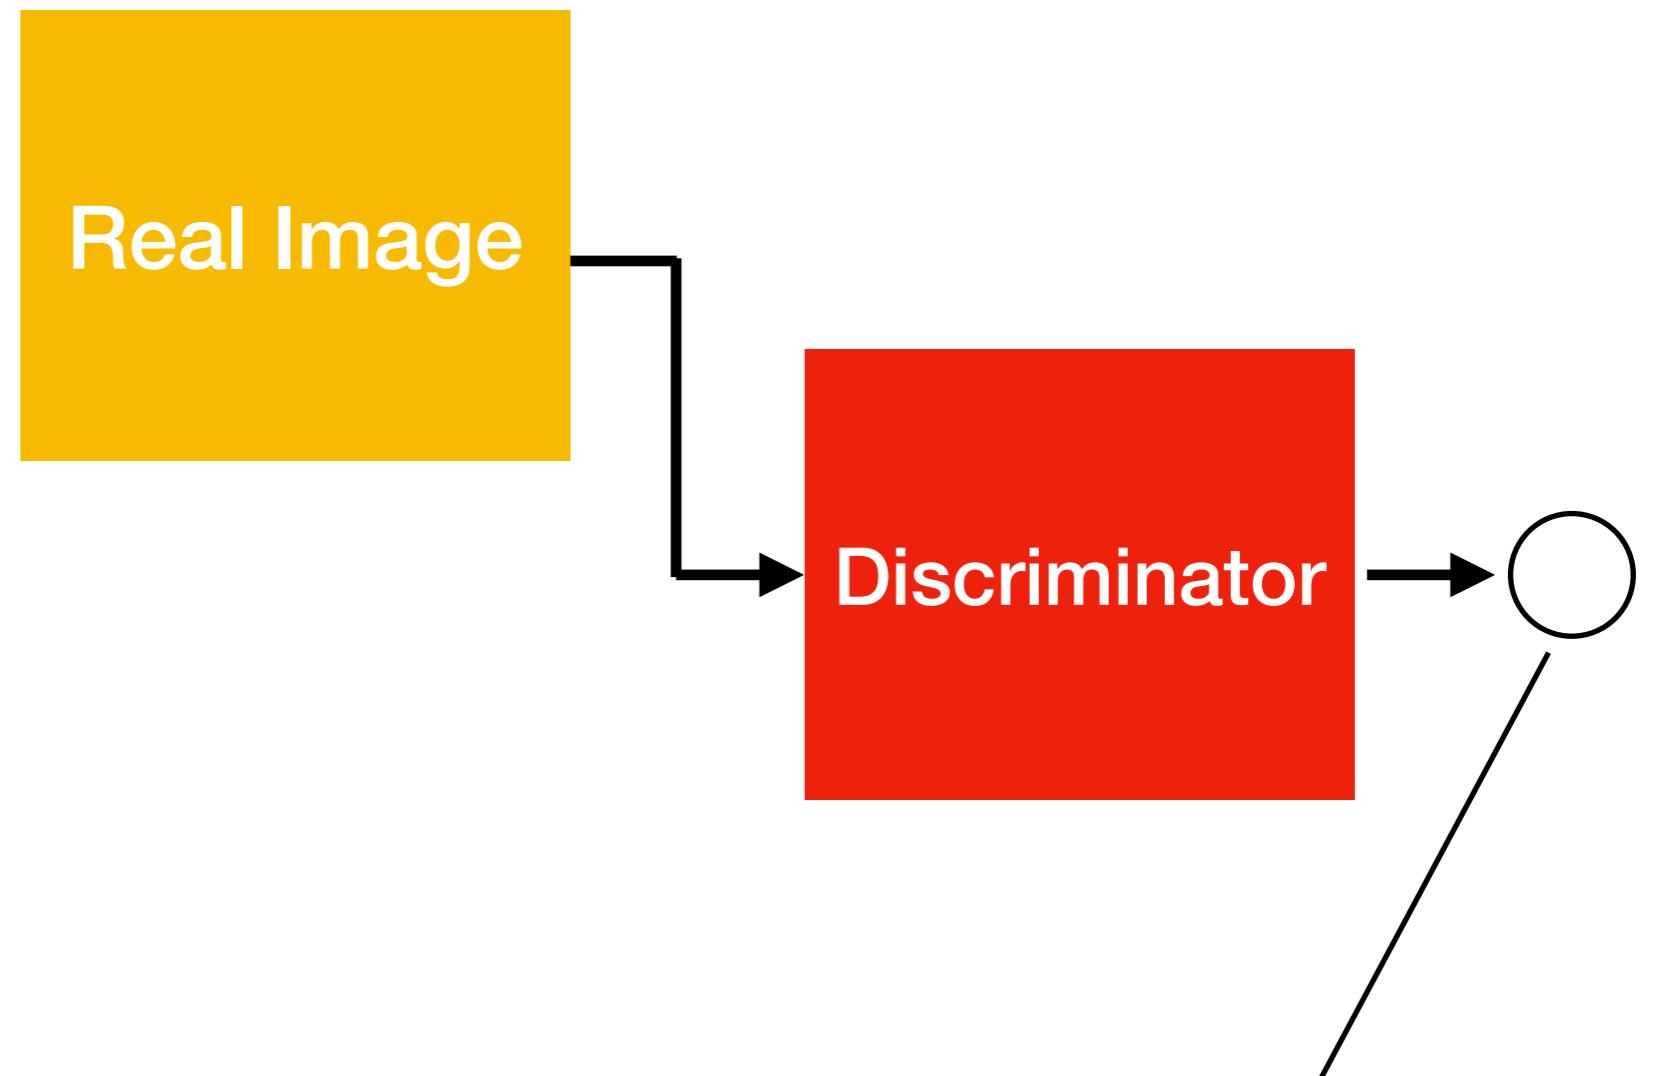
[Submitted on 10 Jun 2014]

Generative Adversarial Networks

[Ian J. Goodfellow](#), [Jean Pouget-Abadie](#), [Mehdi Mirza](#), [Bing Xu](#), [David Warde-Farley](#), [Sherjil Ozair](#), [Aaron Courville](#), [Yoshua Bengio](#)

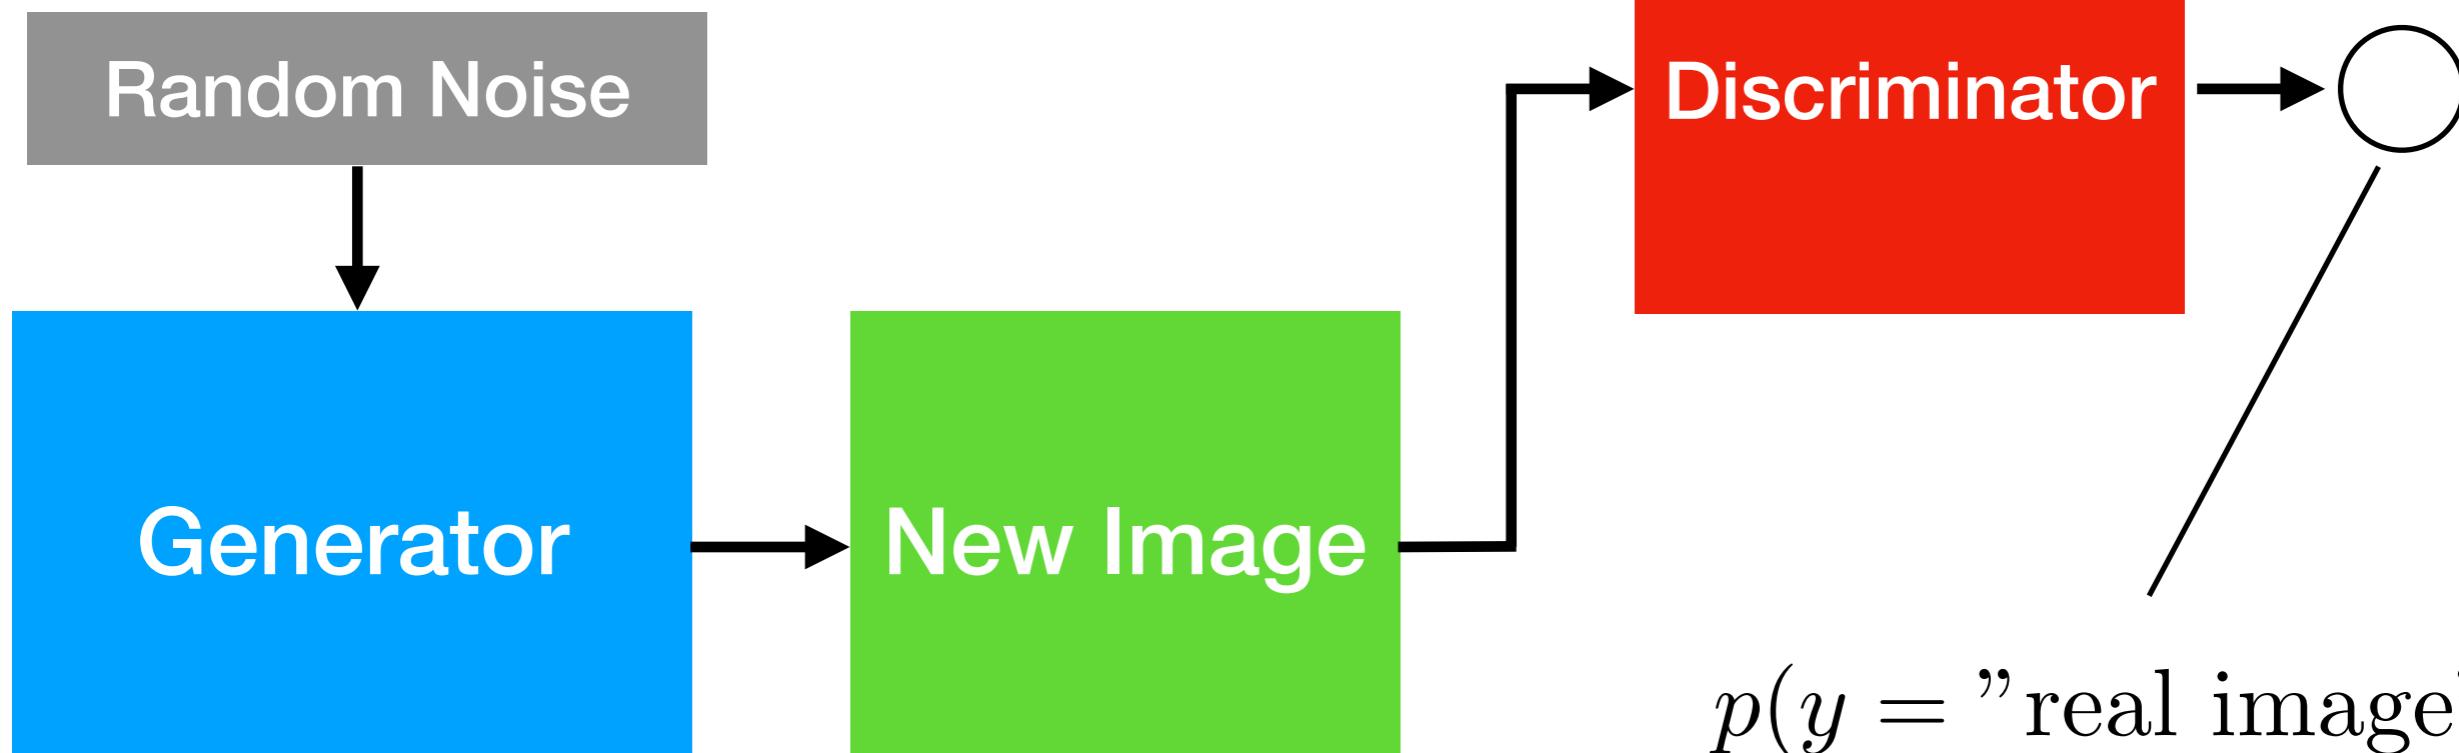
We propose a new framework for estimating generative models via an adversarial process, in which we simultaneously train two models: a generative model G that captures the data distribution, and a discriminative model D that estimates the probability that a sample came from the training data rather than G . The training procedure for G is to maximize the probability of D making a mistake. This framework corresponds to a minimax two-player game. In the space of arbitrary functions G and D , a unique solution exists, with G recovering the training data distribution and D equal to 1/2 everywhere. In the case where G and D are defined by multilayer perceptrons, the entire system can be trained with backpropagation. There is no need for any Markov chains or unrolled approximate inference networks during either training or generation of samples. Experiments demonstrate the potential of the framework through qualitative and quantitative evaluation of the generated samples.

<https://arxiv.org/abs/1406.2661>



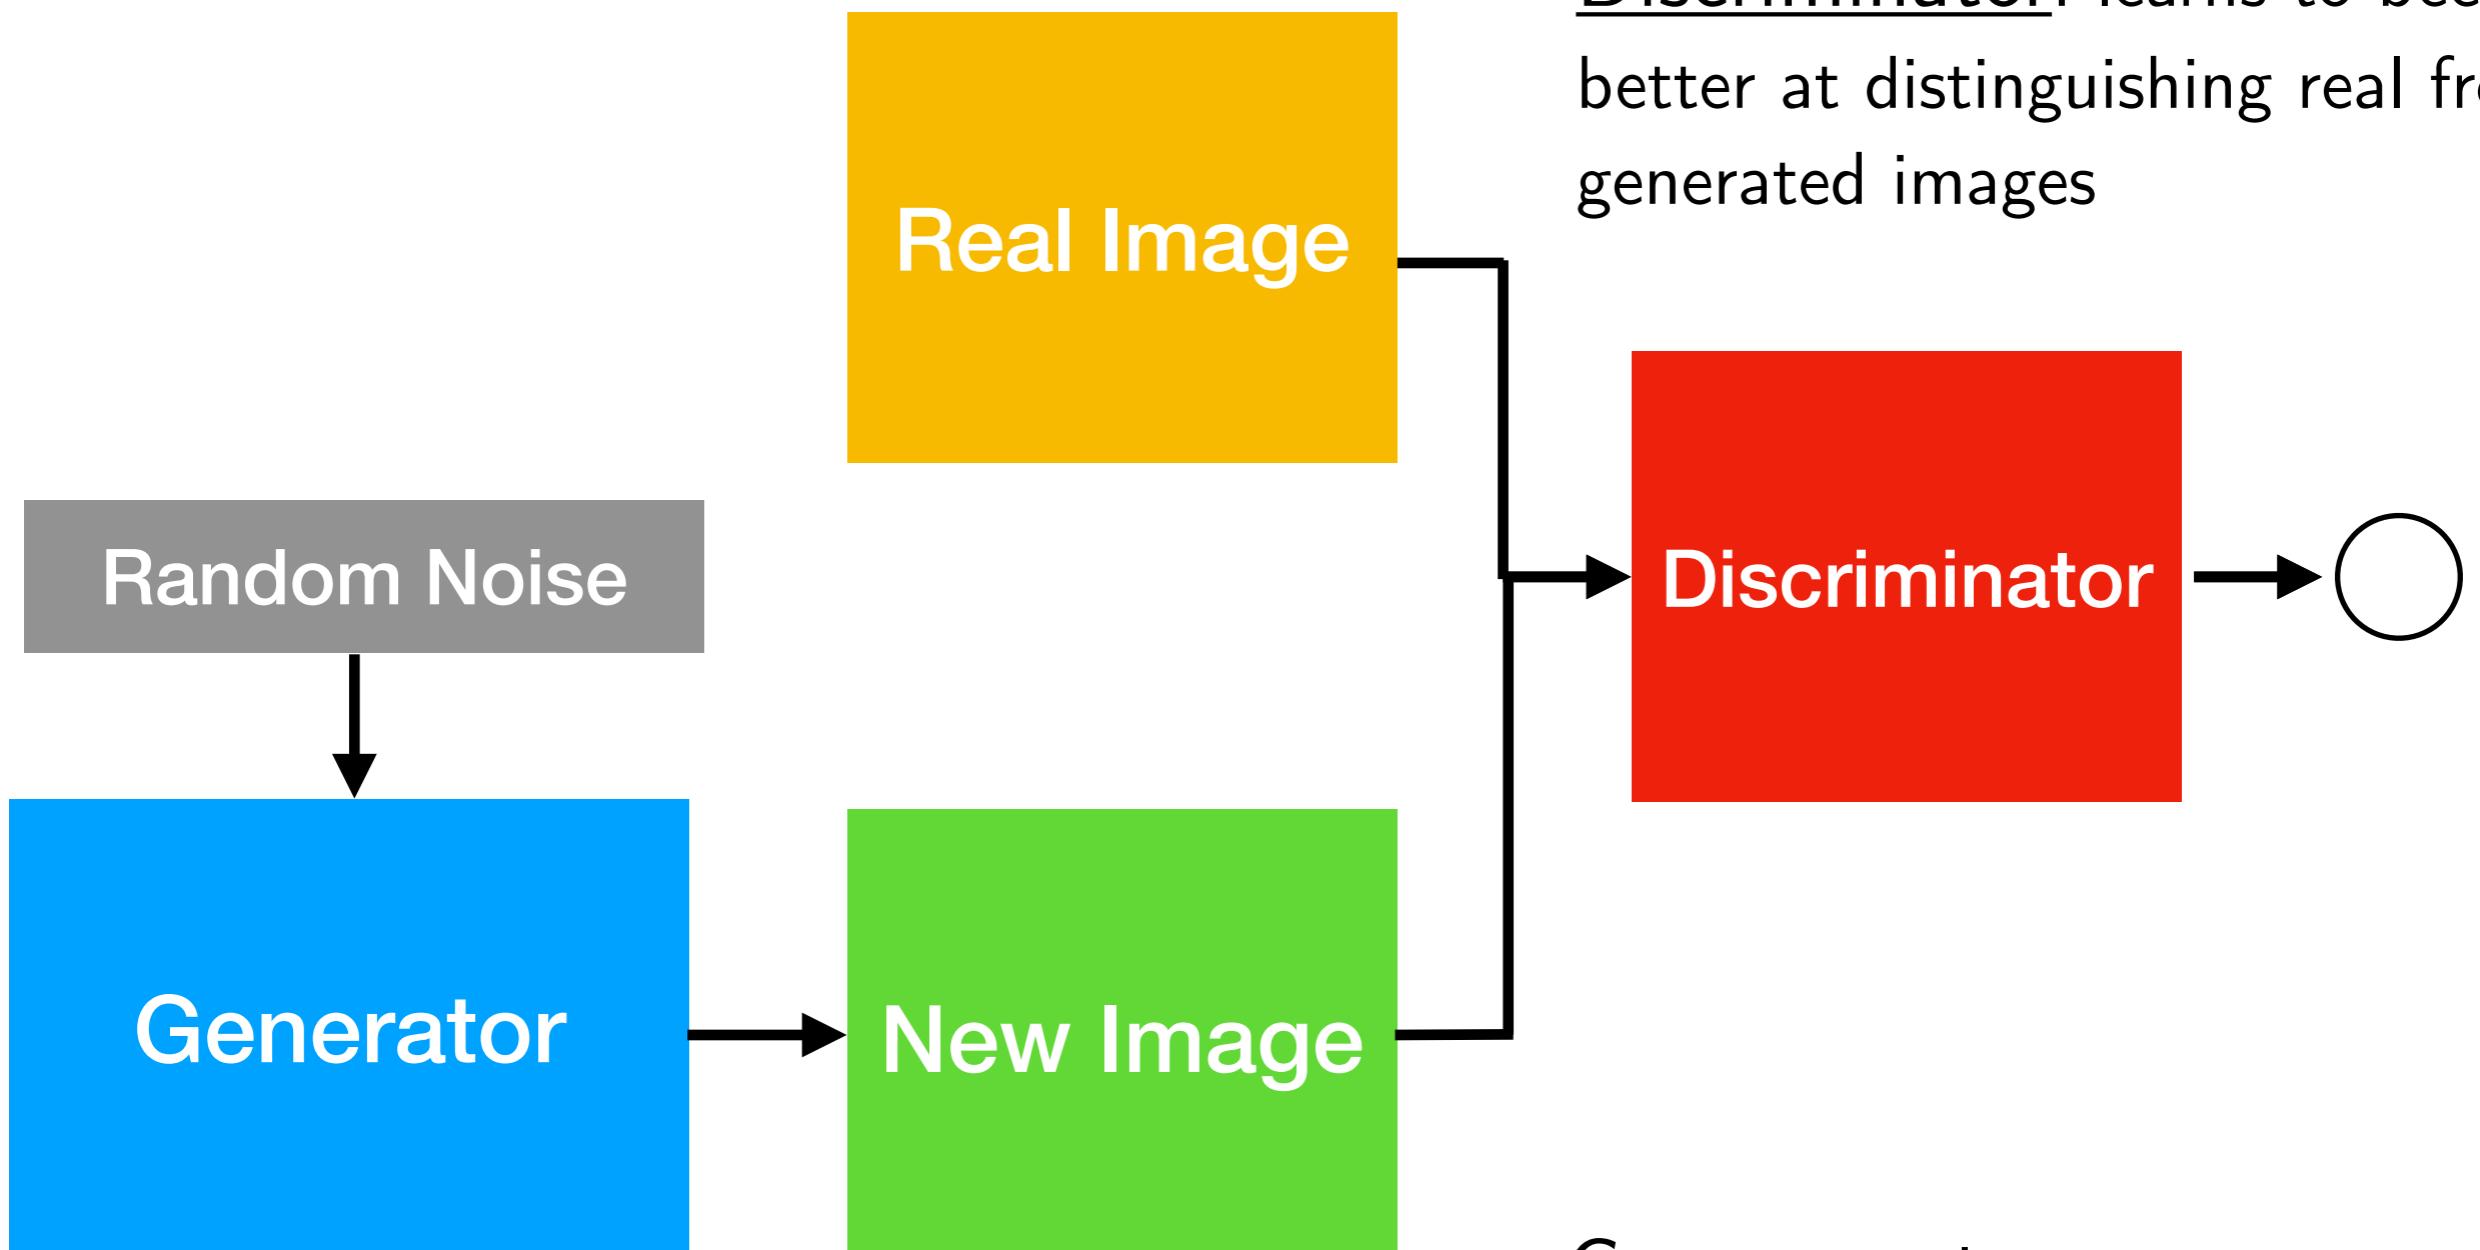
$$p(y = \text{"real image"} | \mathbf{x})$$

$z \sim \mathcal{N}(0, 1)$
or $z \sim U(-1, 1)$



$$p(y = \text{"real image"} | \mathbf{x}')$$

Adversarial Game



Discriminator: learns to become better at distinguishing real from generated images

Generator: learns to generate better images to fool the discriminator

GAN or GANs? Singular vs plural.

If you wonder how the term "generative" fits in the context of your other statistics classes ...

Lipton, Z. C., & Steinhardt, J. (2018). [Troubling Trends in Machine Learning Scholarship](#). arXiv preprint arXiv:1807.03341.

As another example, *generative models* are traditionally models of either the input distribution $p(x)$ or the joint distribution $p(x, y)$. In contrast, discriminative models address the conditional distribution $p(y | x)$ of the label given the inputs. However, in recent works, “generative model” imprecisely refers to any model that produces realistic-looking structured data. On the surface, this may seem consistent with the $p(x)$ definition, but it obscures several shortcomings—for instance, the inability of GANs or VAEs to perform conditional inference (e.g. sampling from $p(x_2 | x_1)$ where x_1 and x_2 are two distinct input features). Bending the term further, some discriminative models are now referred to as generative models on account of producing structured outputs [76], a mistake that we (ZL) make in [47]. Seeking to resolve the confusion and provide historical context, [58] distinguishes between *prescribed* and *implicit* generative models.

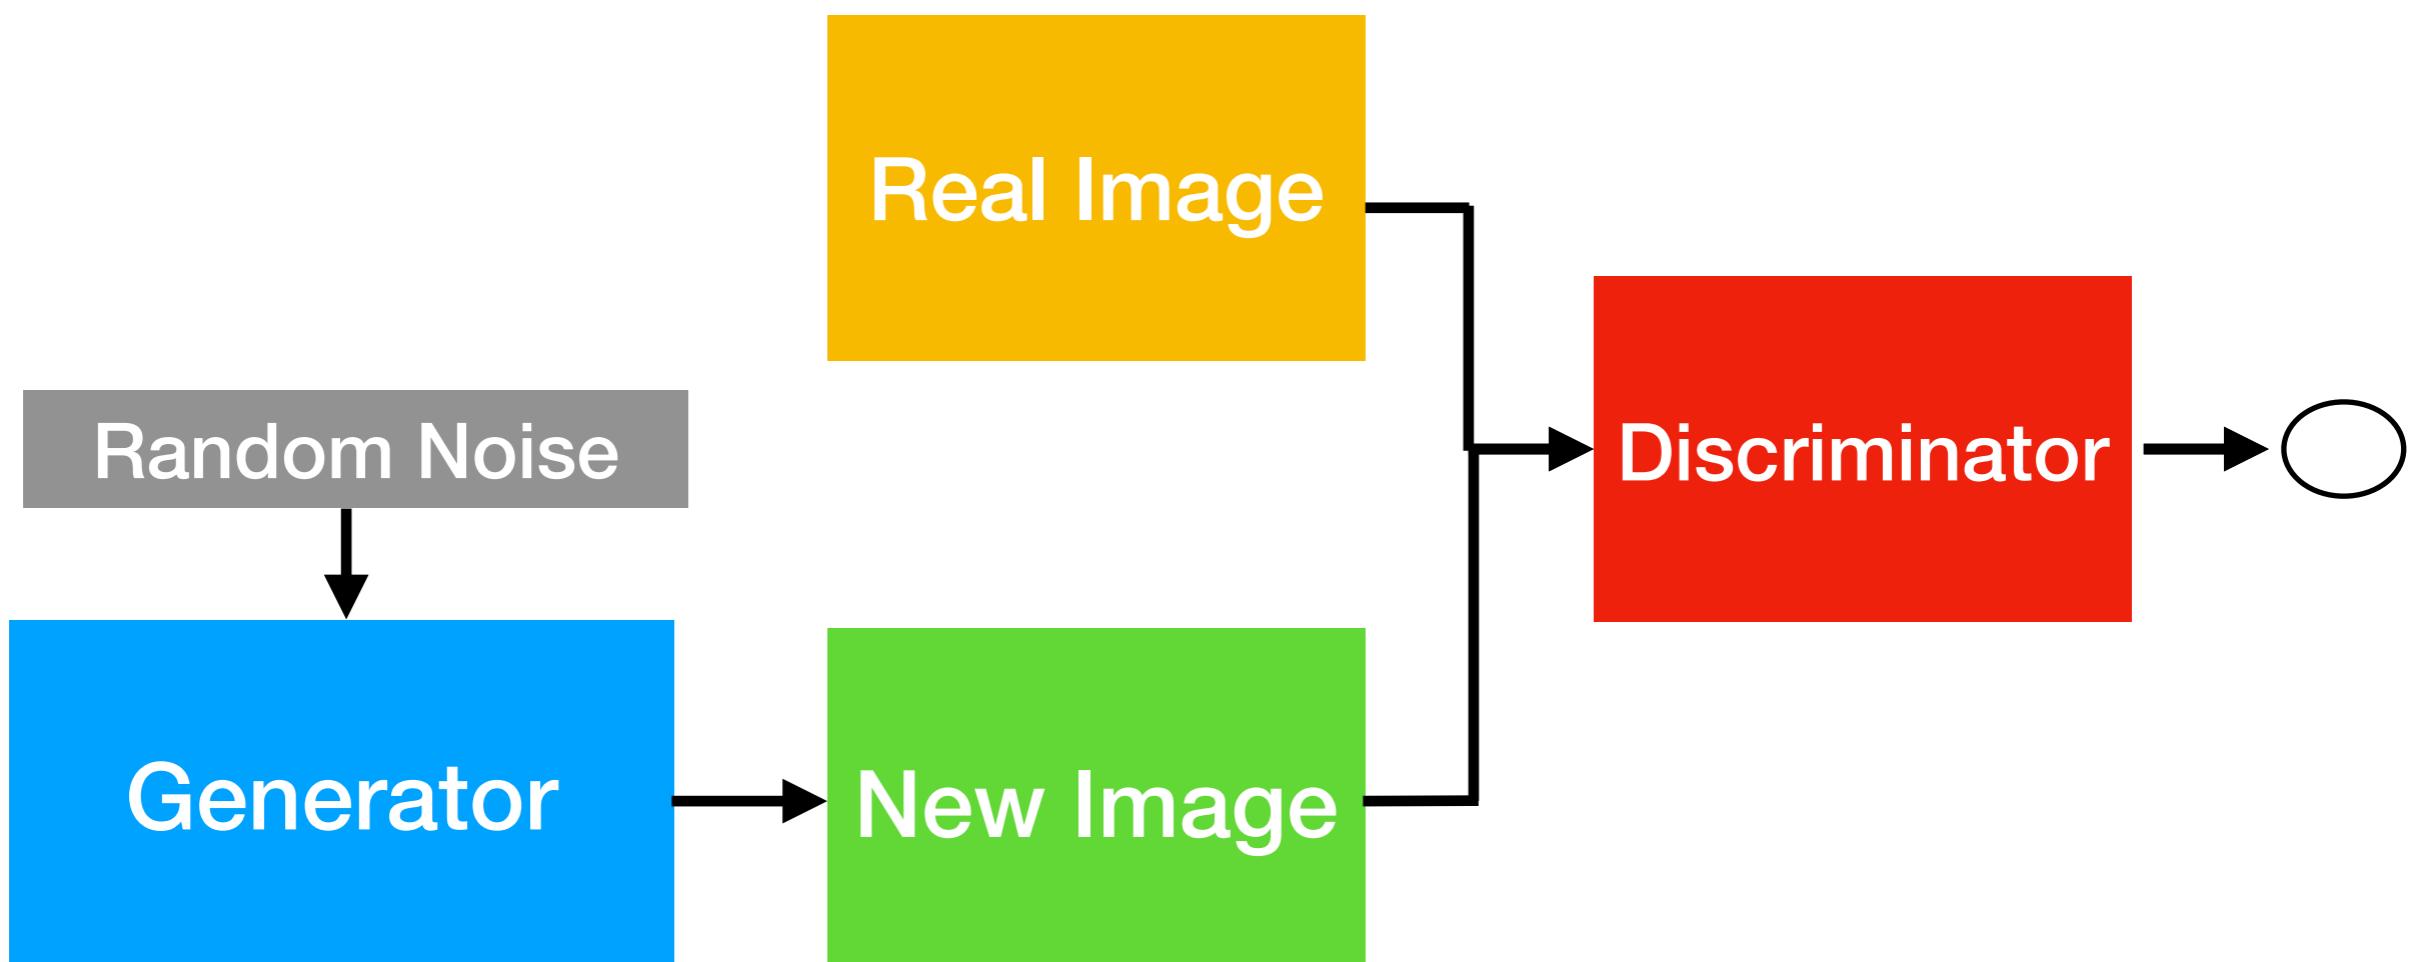
Why Are GANs Are Called Generative Models?

- The generative part comes from the fact that the model "generates" new data
- Usually, generative models use an approximation to compute the usually intractable distribution; here, the discriminator part does that approximation
- So, it does learn $p(x)$
- Vanilla GANs cannot do conditional inference, though

Lecture Overview

1. The Main Idea Behind GAN
- 2. The GAN Training**
3. GANs in Practice
4. GAN Resources

When Does a GAN Converge?



GAN Objective

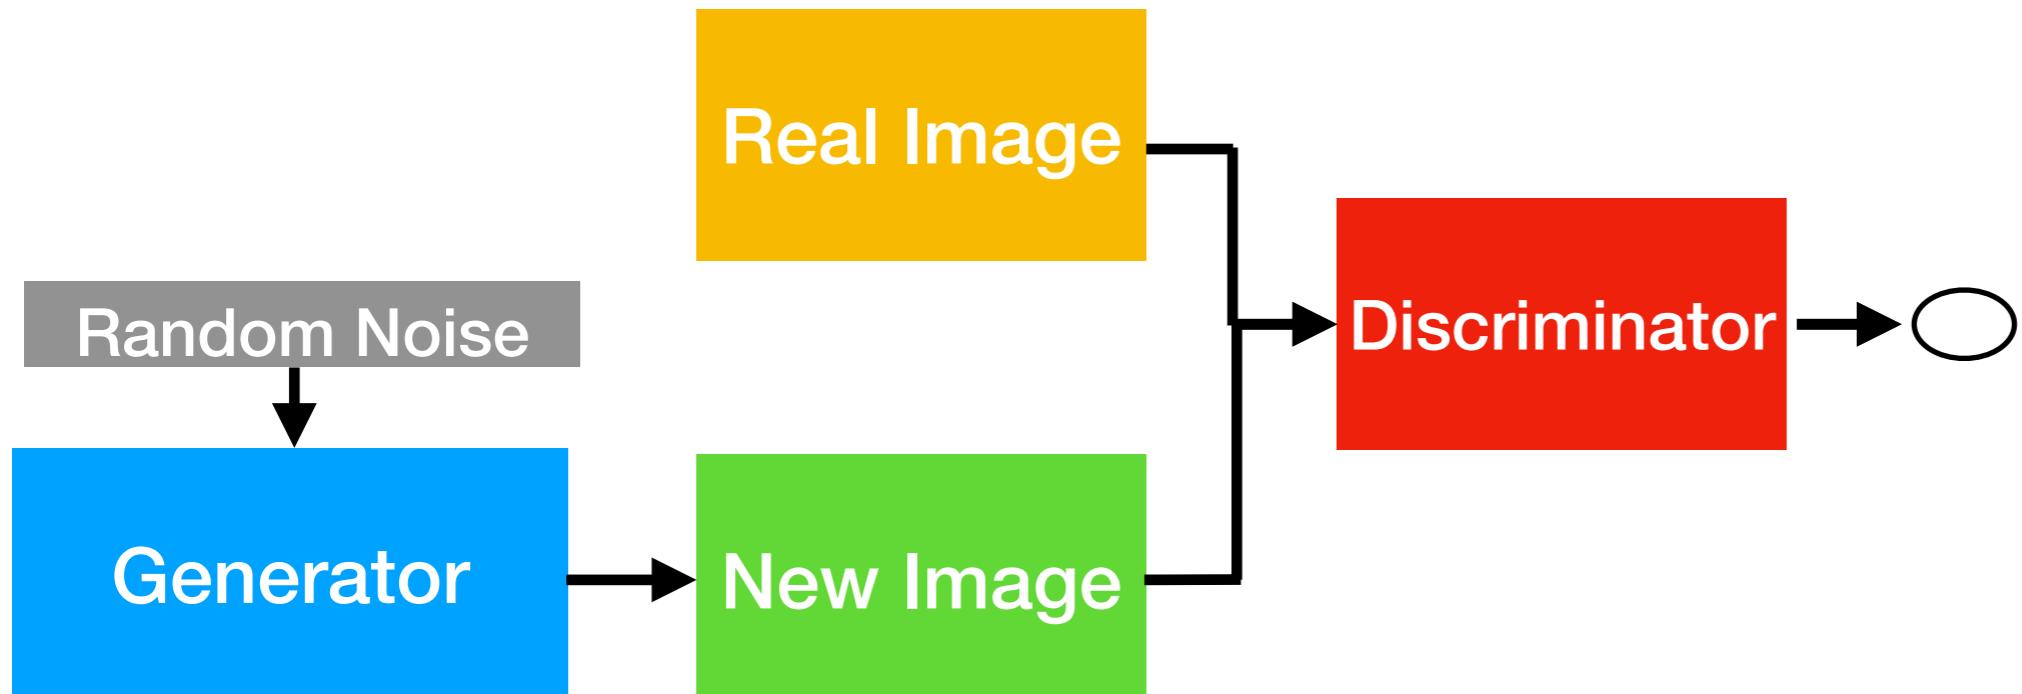
$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

Discriminator gradient for update (gradient ascent):

predict well on real images => probability close to 1	predict well on fake images => probability close to 0
--	--

$$\nabla_{\mathbf{W}_D} \frac{1}{n} \sum_{i=1}^n \left[\log D(\mathbf{x}^{(i)}) + \log (1 - D(G(\mathbf{z}^{(i)}))) \right]$$

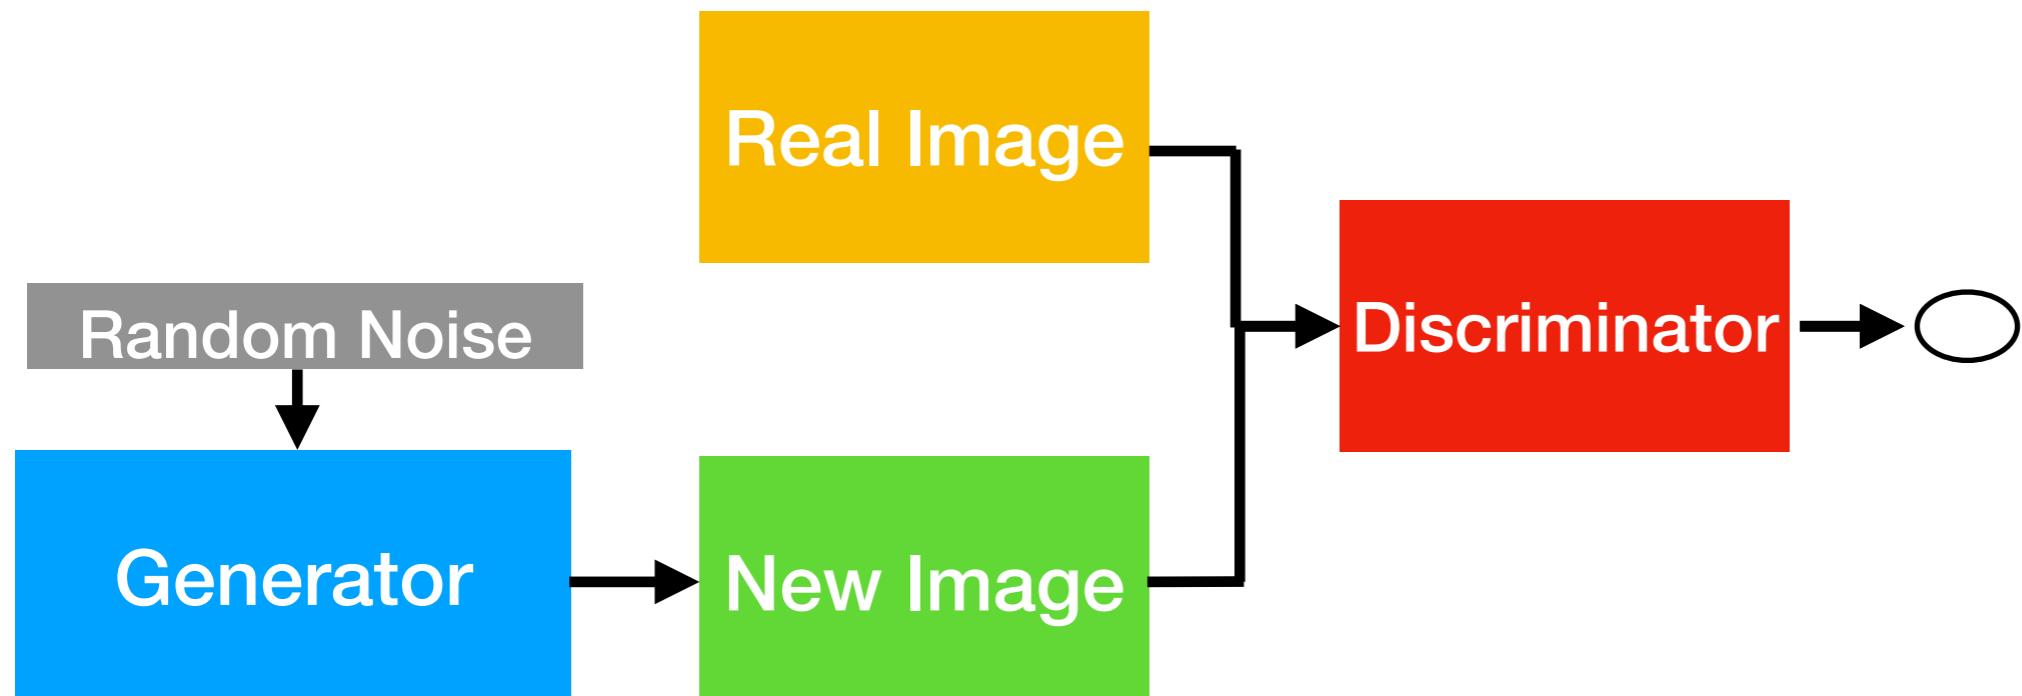


$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

Generator gradient for update (gradient descent):

predict badly on fake images
=> probability close to 1

$$\nabla_{\mathbf{W}_G} \frac{1}{n} \sum_{i=1}^n \log \left(1 - D \left(G \left(\mathbf{z}^{(i)} \right) \right) \right)$$



Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_g(\mathbf{z})$.
- Sample minibatch of m examples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ from data generating distribution $p_{\text{data}}(\mathbf{x})$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(\mathbf{x}^{(i)}) + \log (1 - D(G(\mathbf{z}^{(i)}))) \right].$$

end for

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_g(\mathbf{z})$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(\mathbf{z}^{(i)}))).$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

- Goodfellow, Ian, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. "[Generative Adversarial Nets](#)." In *Advances in Neural Information Processing Systems*, pp. 2672-2680. 2014.

GAN Convergence

- Converges when Nash-equilibrium (Game Theory concept) is reached in the minmax (zero-sum) game

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

- Nash-Equilibrium in Game Theory is reached when the actions of one player won't change depending on the opponent's actions
- Here, this means that the GAN produces realistic images and the discriminator outputs random predictions (probabilities close to 0.5)

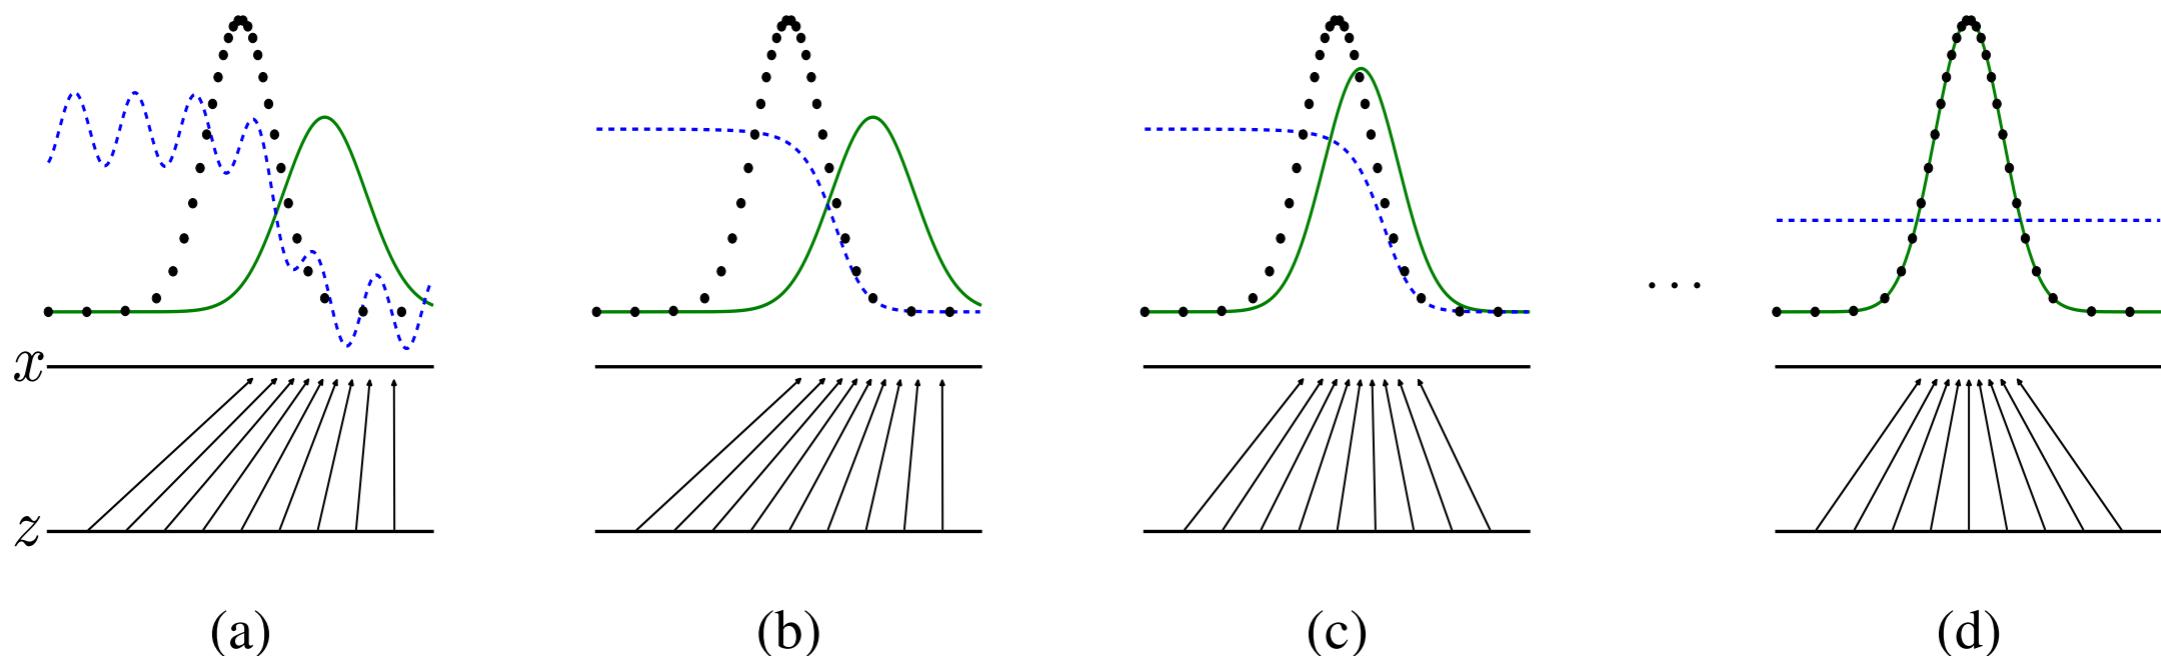


Figure 1: Generative adversarial nets are trained by simultaneously updating the **discriminative** distribution (D , blue, dashed line) so that it discriminates between samples from the data generating distribution (black, dotted line) p_x from those of the **generative** distribution p_g (G) (green, solid line). The lower horizontal line is the domain from which z is sampled, in this case uniformly. The horizontal line above is part of the domain of x . The upward arrows show how the mapping $x = G(z)$ imposes the non-uniform distribution p_g on transformed samples. G contracts in regions of high density and expands in regions of low density of p_g . (a) Consider an adversarial pair near convergence: p_g is similar to p_{data} and D is a partially accurate classifier. (b) In the inner loop of the algorithm D is trained to discriminate samples from data, converging to $D^*(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_g(x)}$. (c) After an update to G , gradient of D has guided $G(z)$ to flow to regions that are more likely to be classified as data. (d) After several steps of training, if G and D have enough capacity, they will reach a point at which both cannot improve because $p_g = p_{\text{data}}$. The discriminator is unable to differentiate between the two distributions, i.e. $D(x) = \frac{1}{2}$.

- Goodfellow, Ian, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. "[Generative Adversarial Nets](#)." In *Advances in Neural Information Processing Systems*, pp. 2672-2680. 2014.

GAN Training Problems

- Oscillation between generator and discriminator loss
- Mode collapse (generator produces examples of a particular kind only)
- Discriminator is too strong, such that the gradient for the generator vanishes and the generator can't keep up
- Discriminator is too weak, and the generator produces non-realistic images that fool it too easily (rare problem, though)

GAN Training Problems

- Discriminator is too strong, such that the gradient for the generator vanishes and the generator can't keep up
- Can be fixed as follows:

Instead of gradient descent with

$$\nabla_{\mathbf{W}_G} \frac{1}{n} \sum_{i=1}^n \log \left(1 - D \left(G \left(\mathbf{z}^{(i)} \right) \right) \right)$$

Do gradient ascent with

$$\nabla_{\mathbf{W}_G} \frac{1}{n} \sum_{i=1}^n \log \left(D \left(G \left(\mathbf{z}^{(i)} \right) \right) \right)$$

GAN Loss Function in Practice

(will be more clear in the code examples)

Discriminator

- Maximize prediction probability of classifying real as real and fake as fake
- Remember maximizing log likelihood is the same as minimizing negative log likelihood (i.e., minimizing cross-entropy)

Generator

- Minimize likelihood of the discriminator to make correct predictions (predict fake as fake; real as real), which can be achieved by maximizing the cross-entropy
- This doesn't work well in practice though because of gradient issues (zero gradient if the discriminator makes correct predictions, which is not what we want for the generator)
- Better: flip labels and minimize cross entropy (force the discriminator to output high probability for fake if an image is real, and high probability for real if an image is fake)

Lecture Overview

1. The Main Idea Behind GAN
2. The GAN Training
- 3. GANs in Practice**
4. GAN Resources

GANs In Practice (1)

Set up a vector of 1's for the real images and a vector of 0's for the fake images

```
valid = torch.ones(targets.size(0)).float().to(device)
fake = torch.zeros(targets.size(0)).float().to(device)
```

GANs In Practice (2)

Set up a vector of 1's for the real images and a vector of 0's for the fake images

```
valid = torch.ones(targets.size(0)).float().to(device)
fake = torch.zeros(targets.size(0)).float().to(device)
```

Generate new images from random noise

```
# Make new images
z = torch.zeros((targets.size(0), LATENT_DIM)).uniform_(-1.0, 1.0).to(device)
generated_features = model.generator_forward(z)
```

GANs In Practice (3)

Setup a vector of 1's for the real images and a vector of 0's for the fake images

```
valid = torch.ones(targets.size(0)).float().to(device)
fake = torch.zeros(targets.size(0)).float().to(device)
```

Generate new images from random noise

```
# Make new images
z = torch.zeros((targets.size(0), LATENT_DIM)).uniform_(-1.0, 1.0).to(device)
generated_features = model.generator_forward(z)
```

Generator Loss:

Minimizing likelihood that discriminator makes a correct prediction can be achieved by maximizing likelihood that discriminator makes a wrong prediction (predicting valid images). Maximizing likelihood is the same as minimizing negative log likelihood

```
# Loss for fooling the discriminator
discr_pred = model.discriminator_forward(generated_features.view(targets.size(0), 1, 28, 28))
gener_loss = F.binary_cross_entropy(discr_pred, valid)
```

GANs In Practice (3)

Setup a vector of 1's for the real images and a vector of 0's for the fake images

For reference, from Lecture 8:

$$H_a(y) = - \sum_i \left(y^{[i]} \log(a^{[i]}) + (1 - y^{[i]}) \log(1 - a^{[i]}) \right)$$

Binary Cross Entropy

Generator Loss:

Minimizing likelihood that discriminator makes a correct prediction can be achieved by maximizing likelihood that discriminator makes a wrong prediction (predicting valid images). Maximizing likelihood is the same as minimizing negative log likelihood

```
# Loss for fooling the discriminator
discr_pred = model.discriminator_forward(generated_features.view(targets.size(0), 1, 28, 28))

gener_loss = F.binary_cross_entropy(discr_pred, valid)
```

GANs In Practice (4)

Discriminator Loss:

Train discriminator to recognize real images as real



```
discr_pred_real = model.discriminator_forward(features.view(targets.size(0), 1, 28, 28))  
real_loss = F.binary_cross_entropy(discr_pred_real, valid)
```

```
discr_pred_fake = model.discriminator_forward(generated_features.view(targets.size(0), 1, 28, 28).detach())  
fake_loss = F.binary_cross_entropy(discr_pred_fake, fake)
```

```
discr_loss = 0.5*(real_loss + fake_loss)
```



Train discriminator to recognize generated images as fake

Combine fake & real parts

GANs In Practice (5)

Use separate optimizers for generator and discriminator

```
# Make new images
z = torch.zeros((targets.size(0), LATENT_DIM)).uniform_(-1.0, 1.0).to(device)
generated_features = model.generator_forward(z)

# Loss for fooling the discriminator
discr_pred = model.discriminator_forward(generated_features.view(targets.size(0), 1, 28, 28))

gener_loss = F.binary_cross_entropy(discr_pred, valid)

optim_gener.zero_grad()
gener_loss.backward()
optim_gener.step()

# -----
# Train Discriminator
# -----

discr_pred_real = model.discriminator_forward(features.view(targets.size(0), 1, 28, 28))
real_loss = F.binary_cross_entropy(discr_pred_real, valid)

discr_pred_fake = model.discriminator_forward(generated_features.view(targets.size(0), 1, 28, 28).detach())
fake_loss = F.binary_cross_entropy(discr_pred_fake, fake)

discr_loss = 0.5*(real_loss + fake_loss)

optim_discr.zero_grad()
discr_loss.backward()
optim_discr.step()
```

GANs In Practice (5)

Use separate optimizers for generator and discriminator

```
# Make new images
z = torch.zeros((targets.size(0), LATENT_DIM)).uniform_(-1.0, 1.0).to(device)
generated_features = model.generator_forward(z)

# Loss for fooling the discriminator
discr_pred = model.discriminator_forward(generated_features.view(targets.size(0), 1, 28, 28))

gener_loss = F.binary_cross_entropy(discr_pred, valid)

optim_gener.zero_grad()
gener_loss.backward()
opt
    optim_gener = torch.optim.Adam(model.generator.parameters(), lr=generator_learning_rate)
# -
    optim_discr = torch.optim.Adam(model.discriminator.parameters(), lr=discriminator_learning_rate)
# T
# -----



discr_pred_real = model.discriminator_forward(features.view(targets.size(0), 1, 28, 28))
real_loss = F.binary_cross_entropy(discr_pred_real, valid)

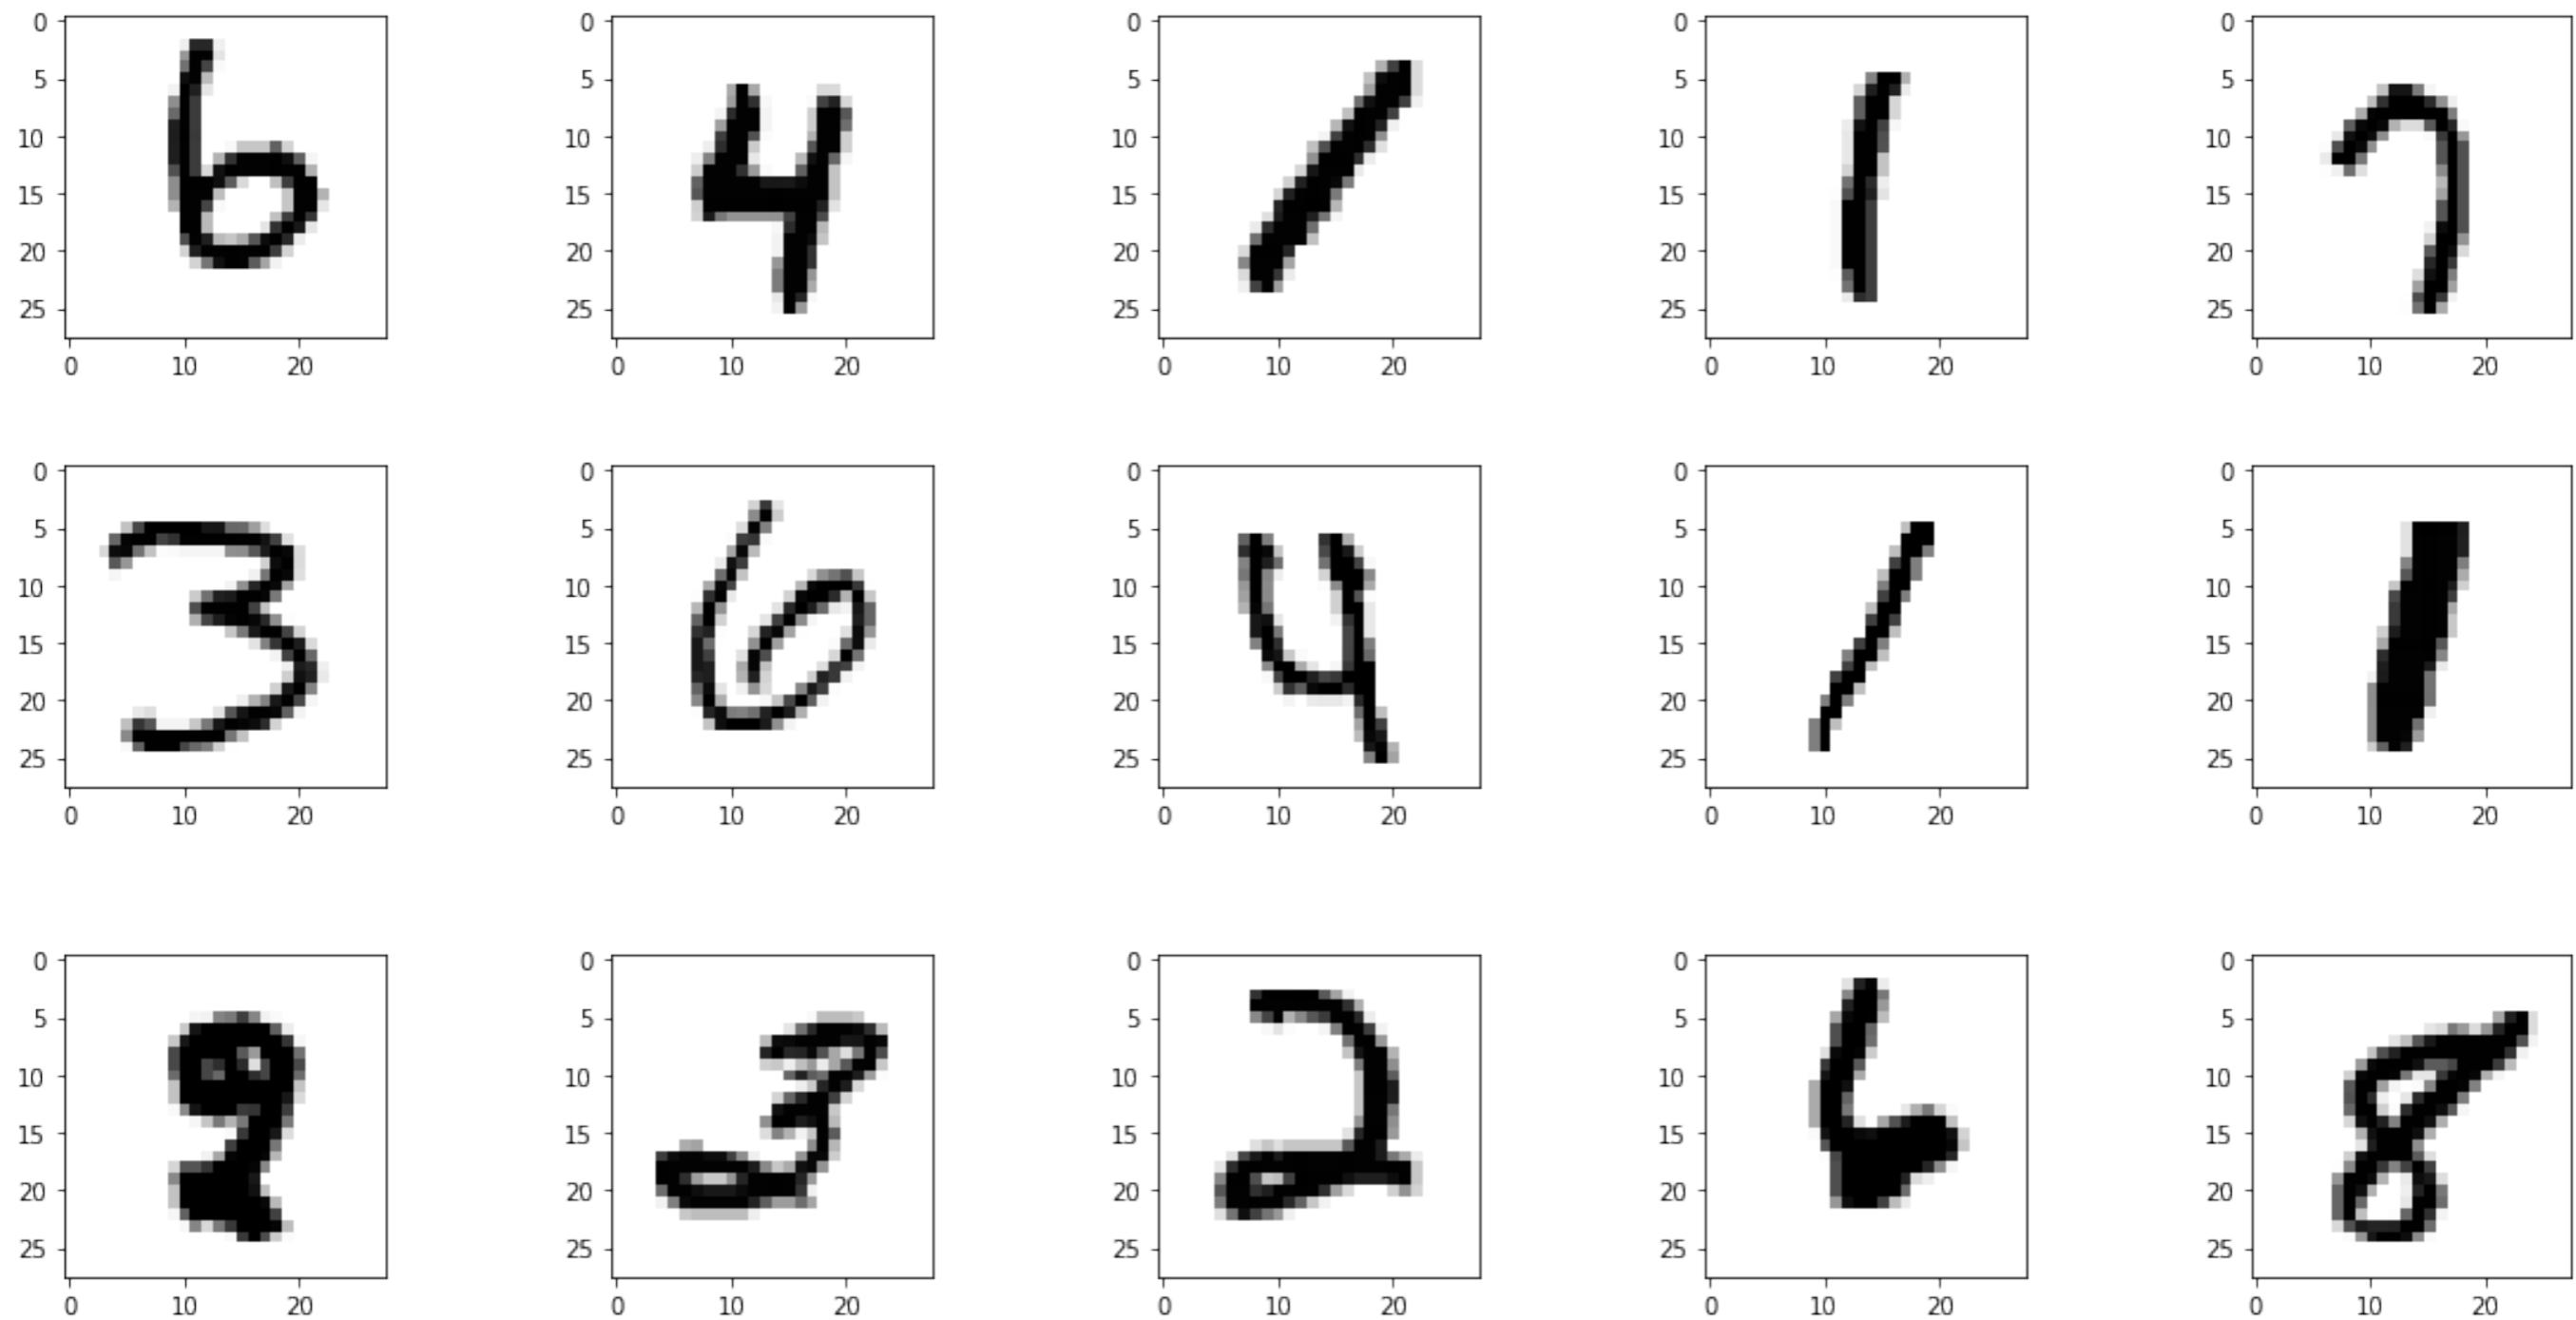
discr_pred_fake = model.discriminator_forward(generated_features.view(targets.size(0), 1, 28, 28).detach())
fake_loss = F.binary_cross_entropy(discr_pred_fake, fake)

discr_loss = 0.5*(real_loss + fake_loss)

optim_discr.zero_grad()
discr_loss.backward()
optim_discr.step()
```

Because we optimize different sets of parameters

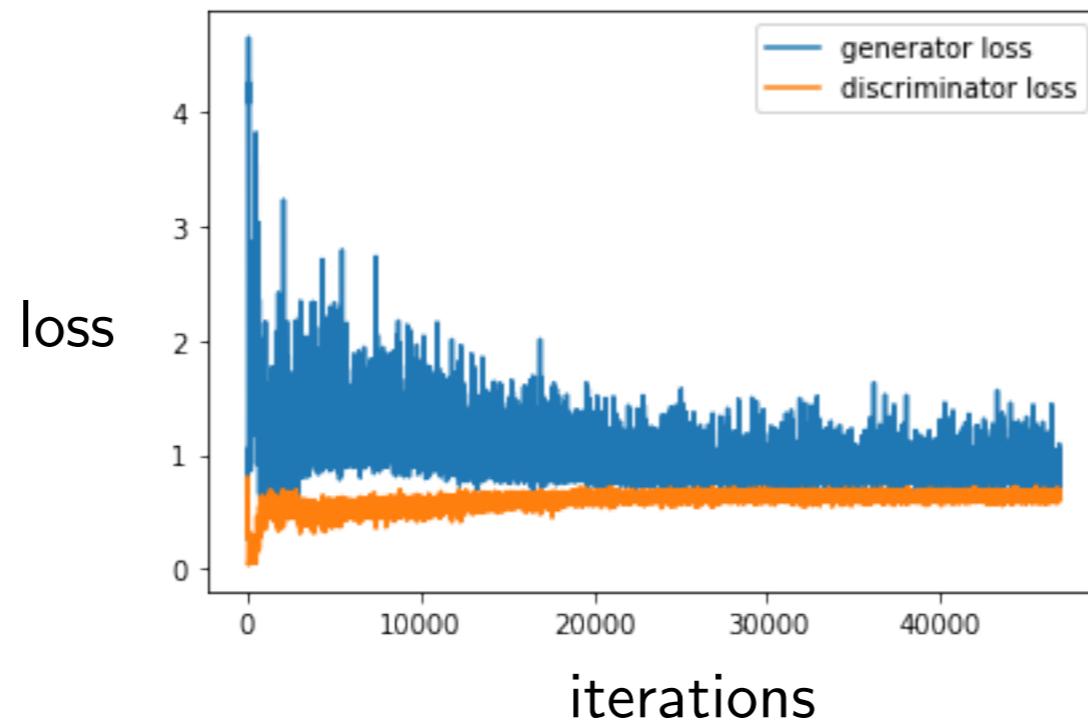
For Reference: Some Real MNIST Images



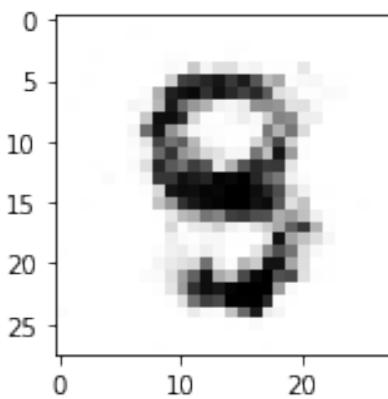
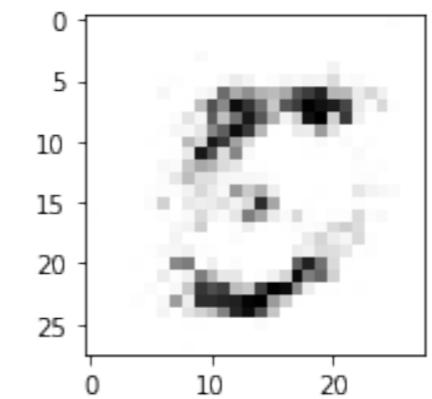
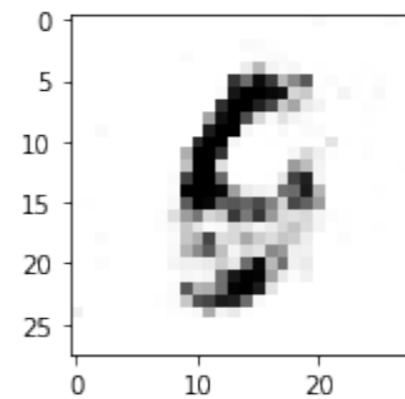
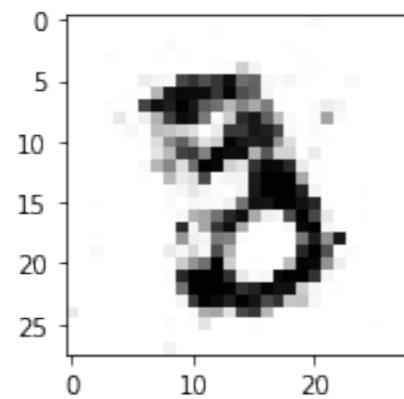
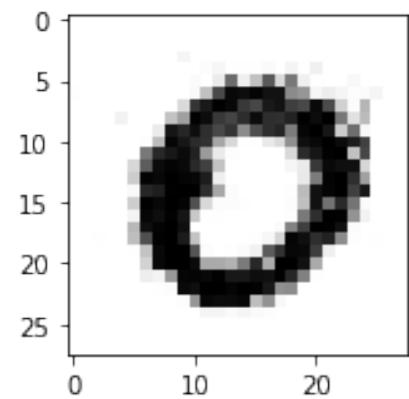
Full Code Examples (1)

MNIST Multilayer Perceptron GAN

https://github.com/rasbt/stat453-deep-learning-ss20/blob/master/L16-gan/code/01_gan-mlp-mnist.ipynb



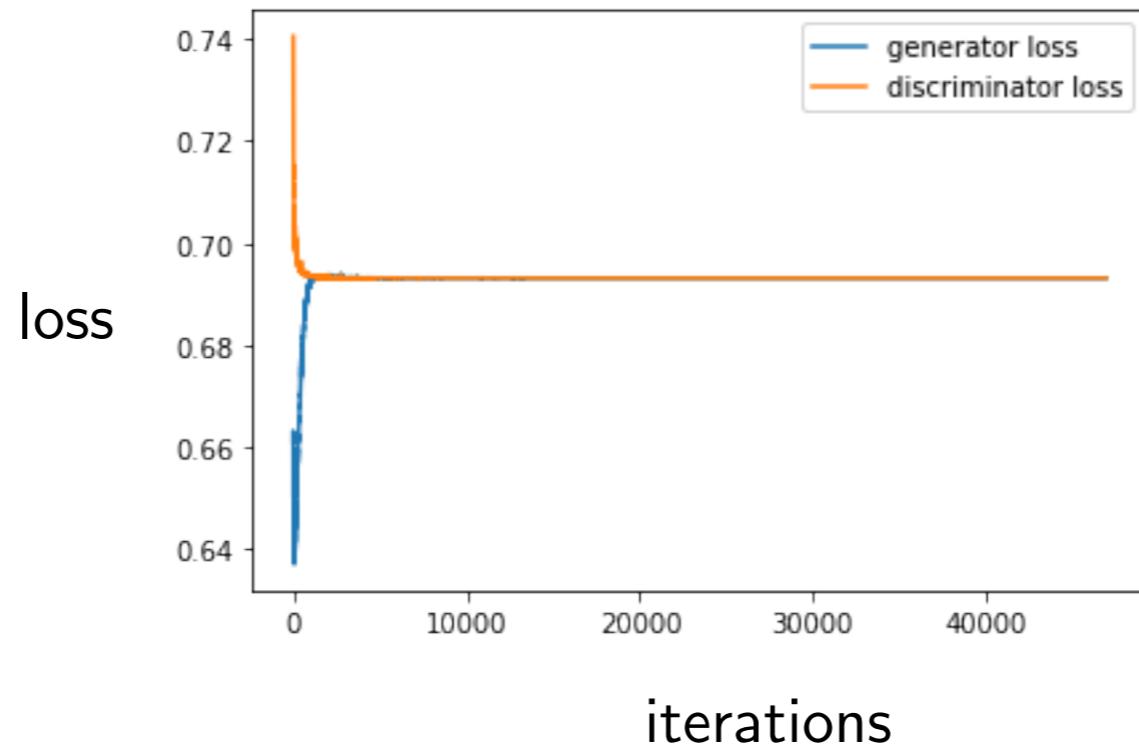
Generated images



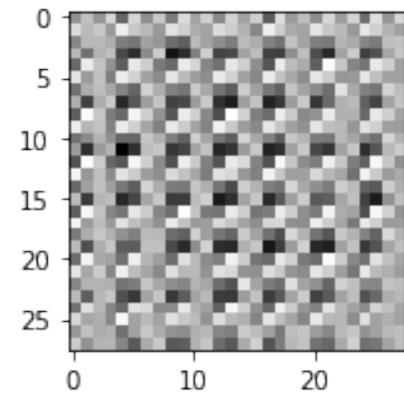
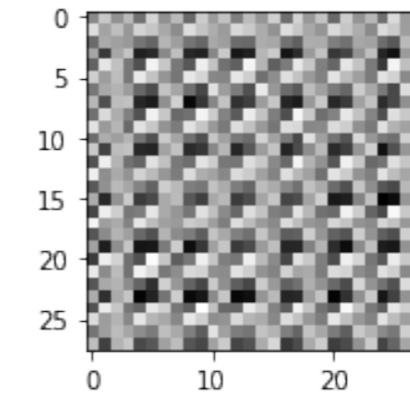
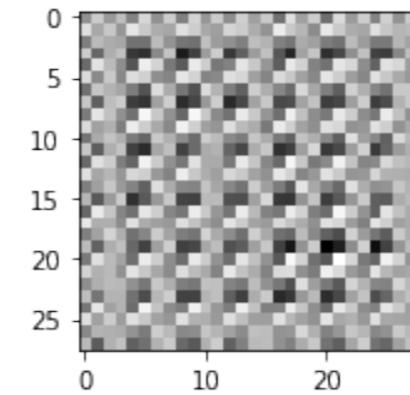
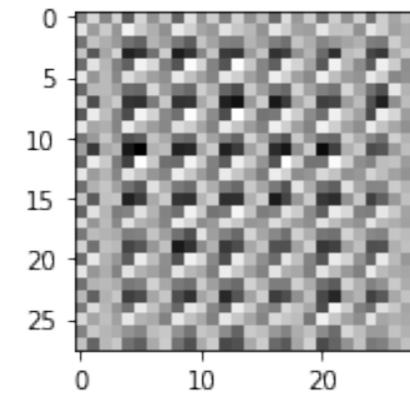
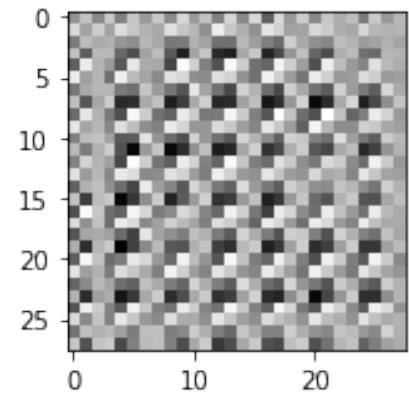
Full Code Examples (2)

MNIST Convolutional GAN that fails

https://github.com/rasbt/stat453-deep-learning-ss20/blob/master/L16-gan/code/02_gan-cnn-mnist-converge-but-fail.ipynb



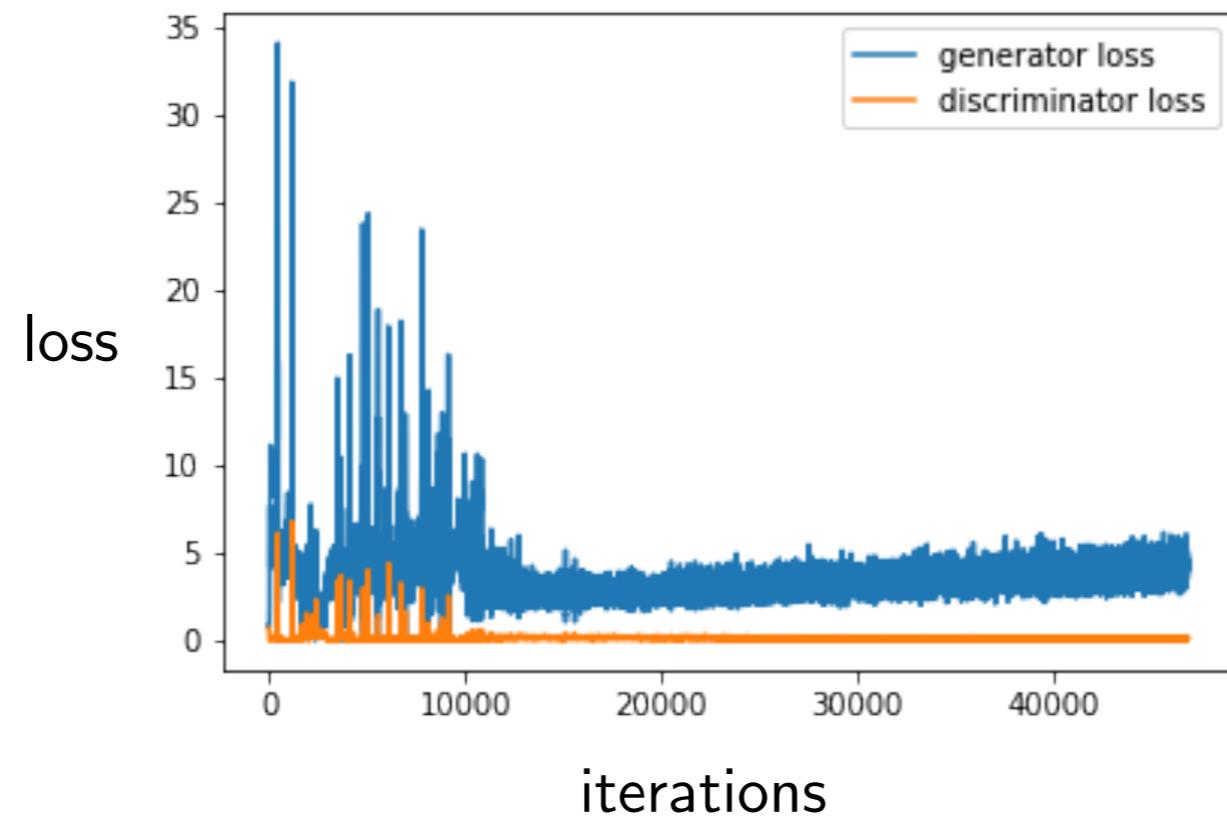
Generated images



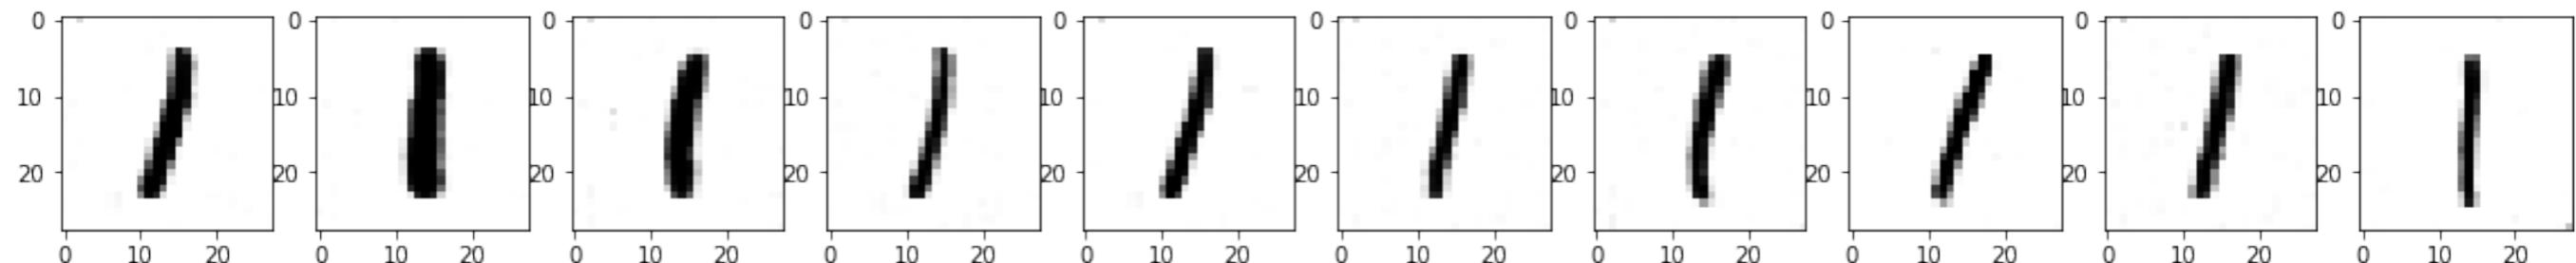
Full Code Examples (3)

MNIST Convolutional GAN with mode collapse

https://github.com/rasbt/stat453-deep-learning-ss20/blob/master/L16-gan/code/03_gan-halfcnn-mnist-mode-collapse.ipynb



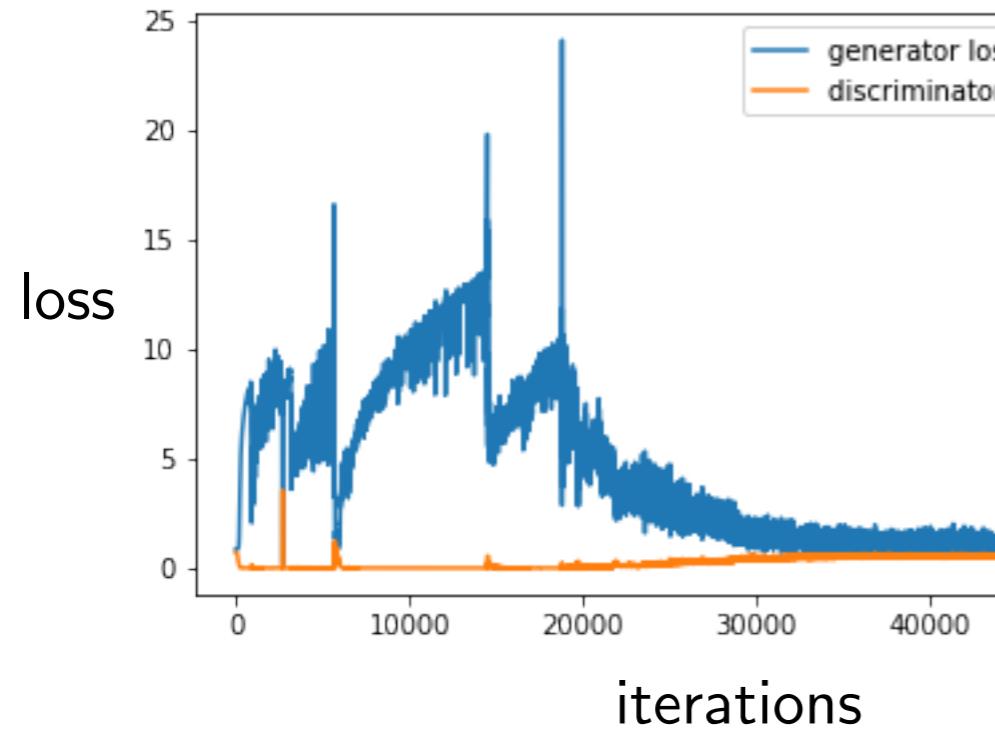
Generated images



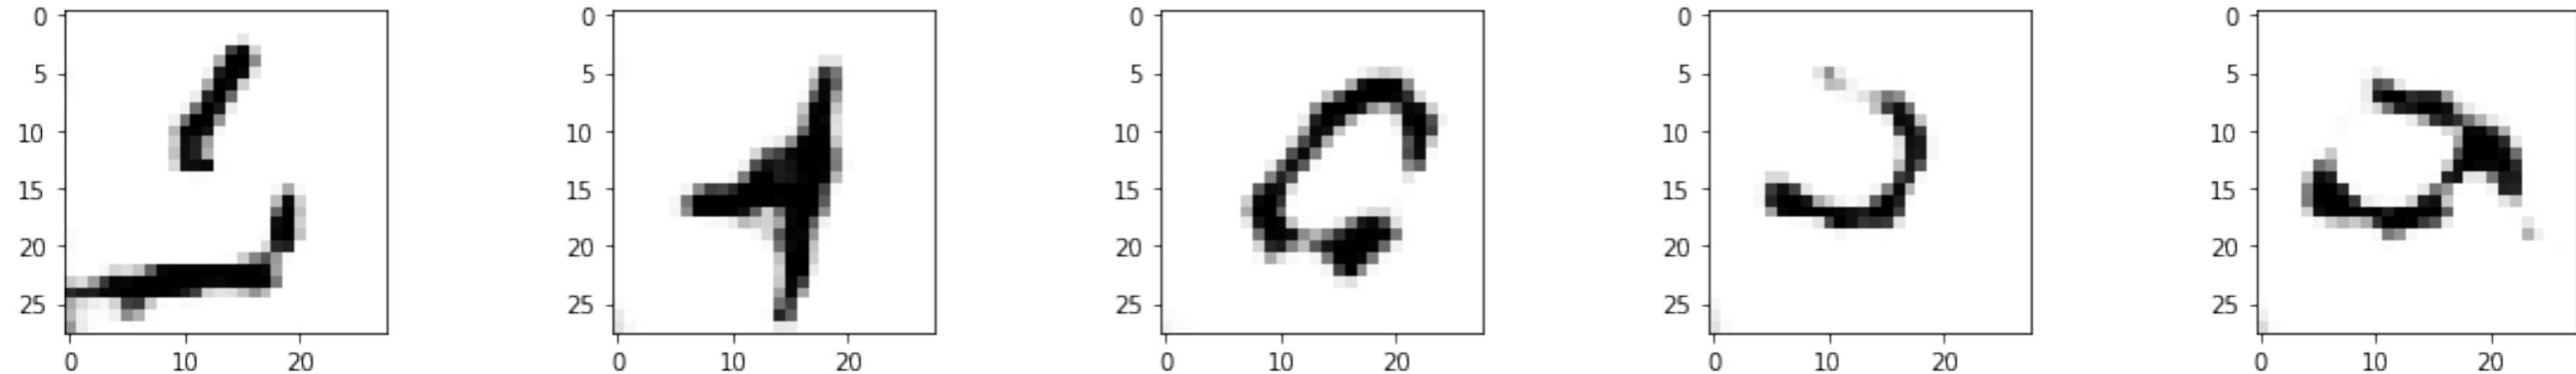
Full Code Examples (4)

MNIST Convolutional GAN -- not great

https://github.com/rasbt/stat453-deep-learning-ss20/blob/master/L16-gan/code/04_gan-cnn-mnist-not-great.ipynb



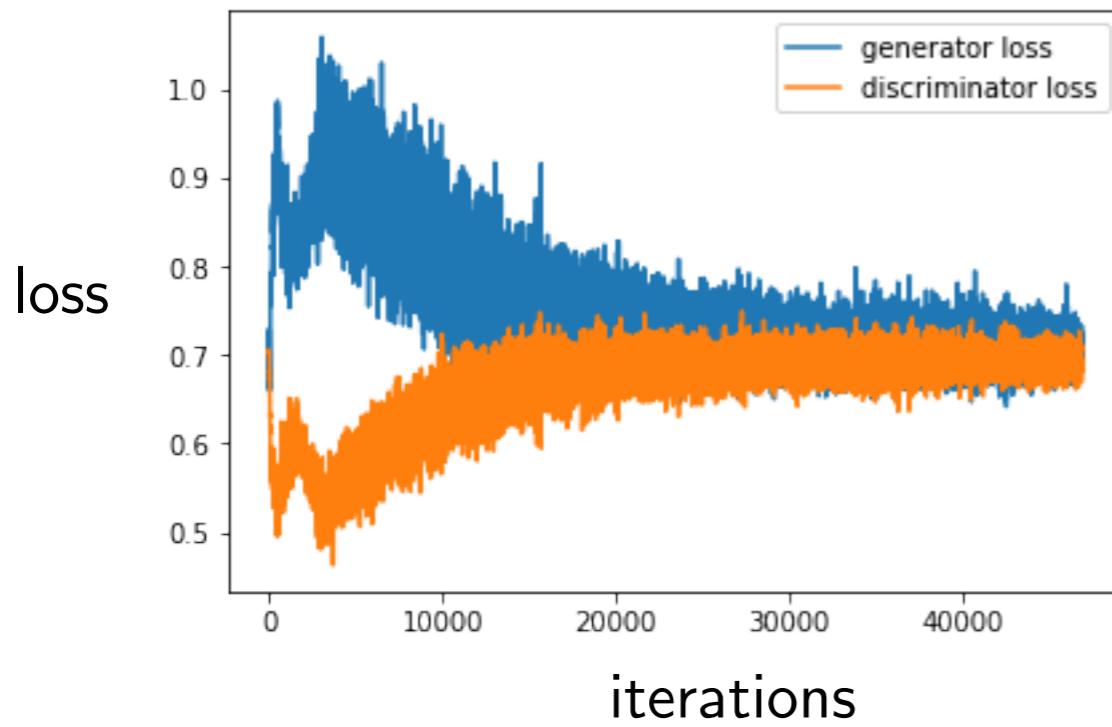
Generated images



Full Code Examples (5)

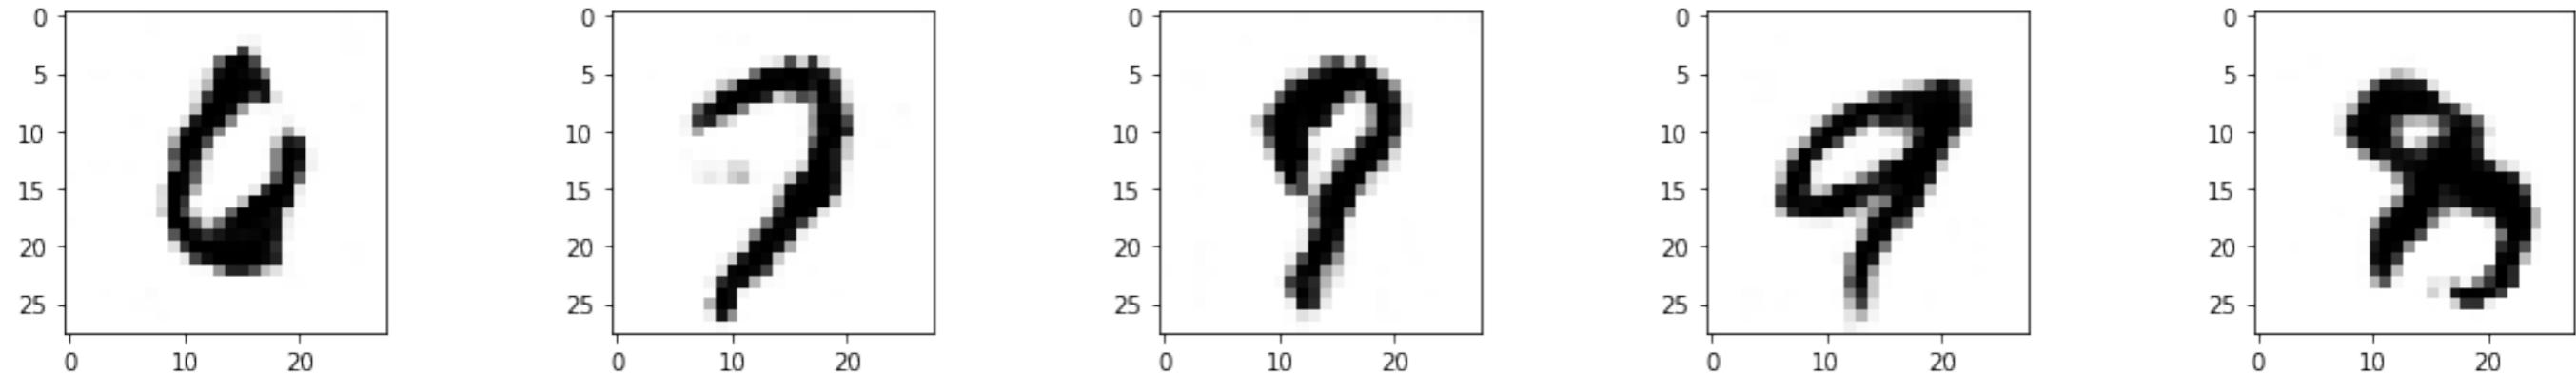
MNIST Convolutional GAN -- relatively good

https://github.com/rasbt/stat453-deep-learning-ss20/blob/master/L16-gan/code/05_gan-cnn-mnist.ipynb



Compared to the previous slide,
this has BatchNorm

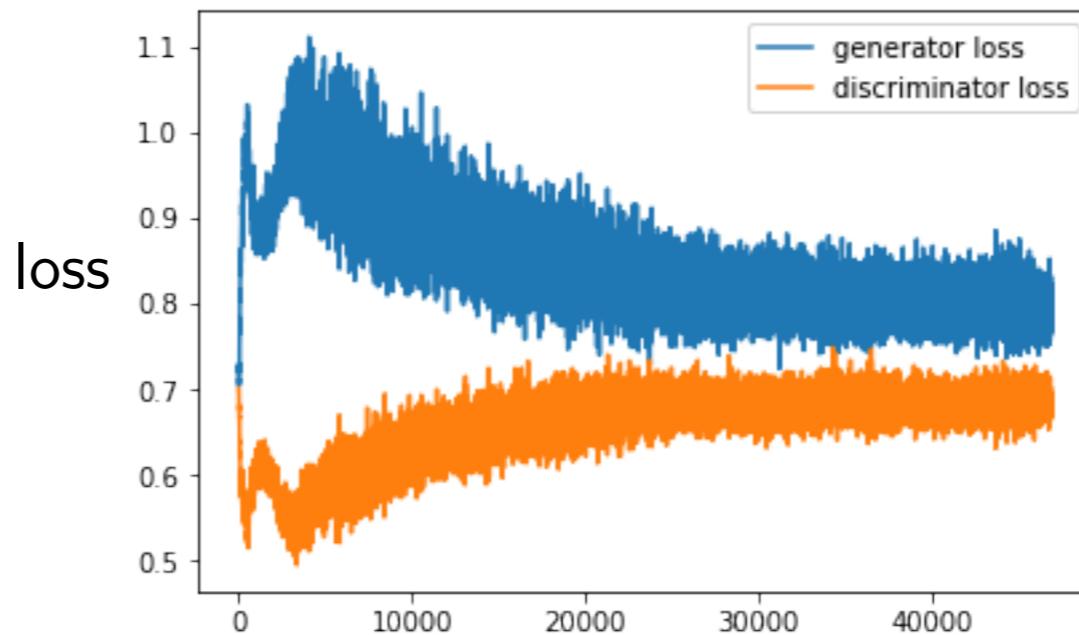
Generated images



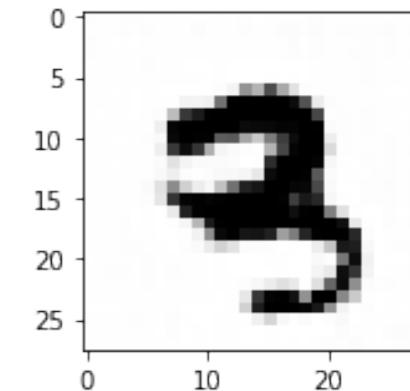
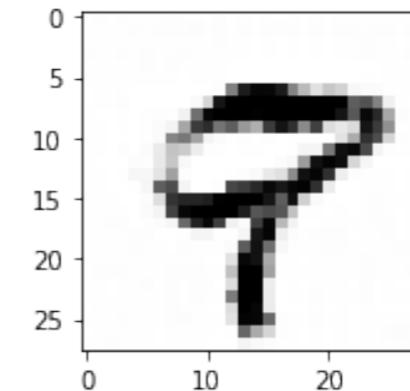
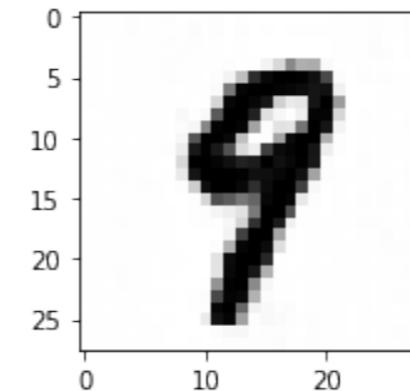
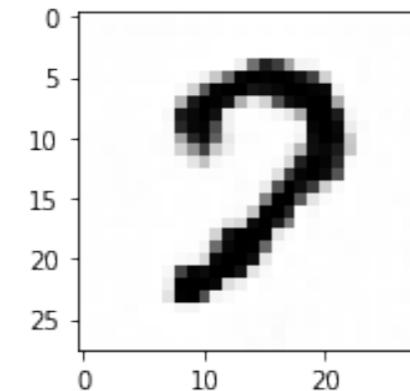
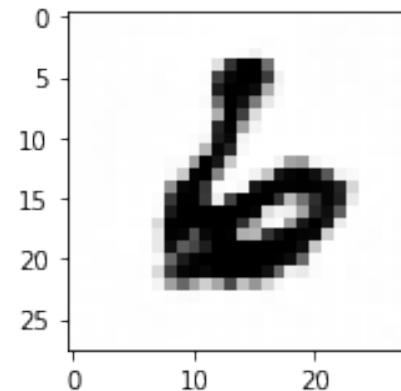
Full Code Examples (6)

Like previous slide but with label smoothing: Replace real images (1's) by 0.9 based on idea in

- Salimans, Tim, Ian Goodfellow, Wojciech Zaremba, Vicki Cheung, Alec Radford, and Xi Chen. "[Improved techniques for training GANs](#)." In *Advances in Neural Information Processing Systems*, pp. 2234-2242. 2016.



https://github.com/rasbt/stat453-deep-learning-ss20/blob/master/L16-gan/code/06_gan-cnn-mnist-label-smoothing.ipynb



Deep Convolutional GAN

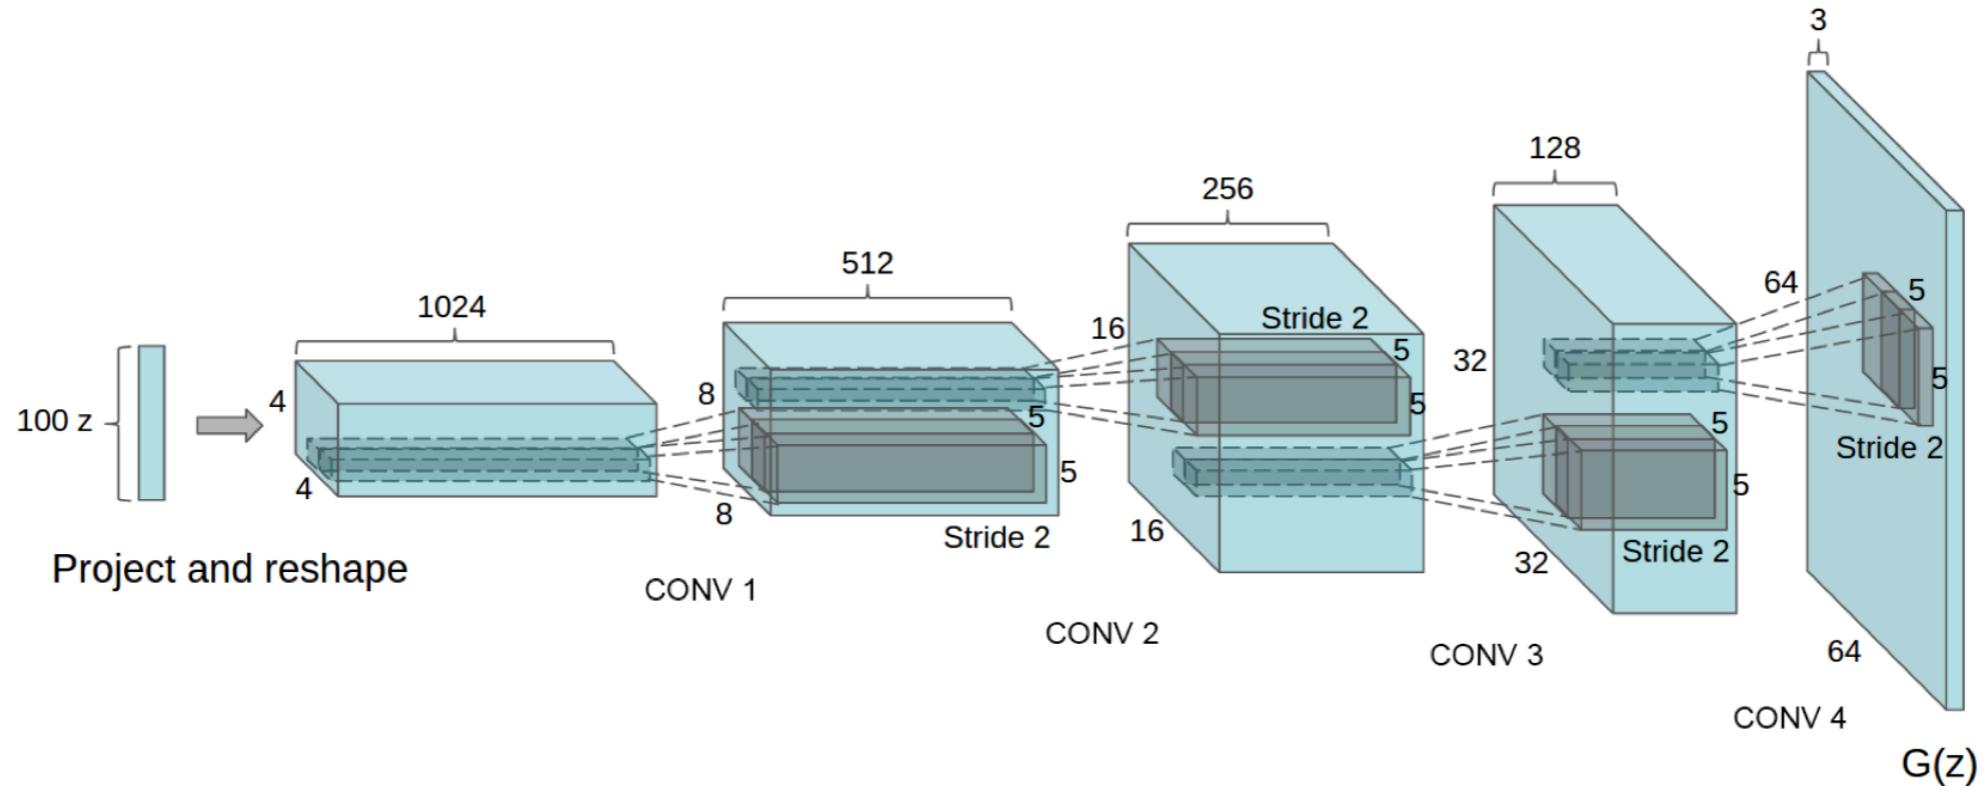
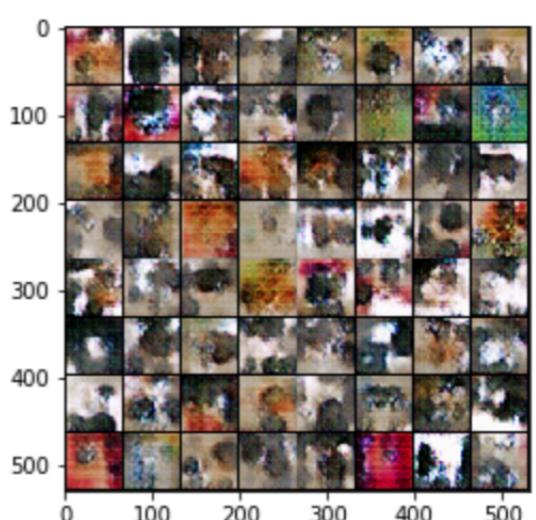
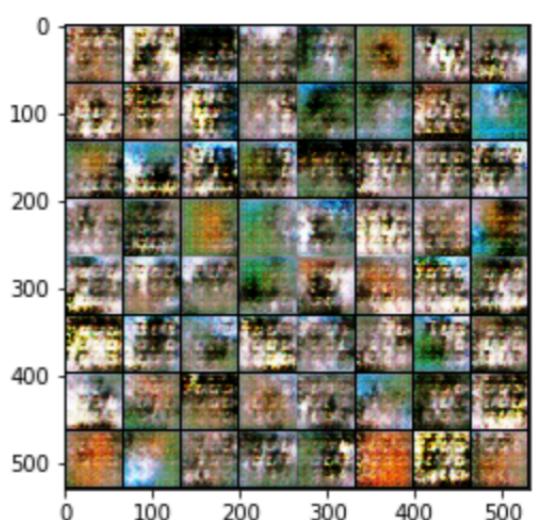
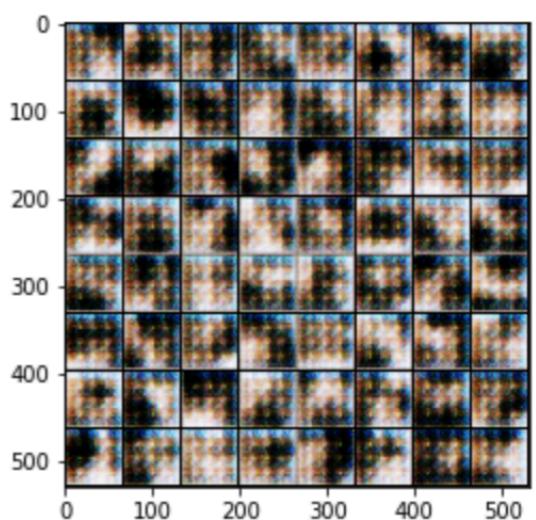


Figure 1: DCGAN generator used for LSUN scene modeling. A 100 dimensional uniform distribution Z is projected to a small spatial extent convolutional representation with many feature maps. A series of four fractionally-strided convolutions (in some recent papers, these are wrongly called deconvolutions) then convert this high level representation into a 64×64 pixel image. Notably, no fully connected or pooling layers are used.

Radford, A., Metz, L., & Chintala, S. (2015). [Unsupervised representation learning with deep convolutional generative adversarial networks](#). arXiv preprint arXiv:1511.06434.

Cats vs Dogs

Training Images



https://github.com/rasbt/stat453-deep-learning-ss20/blob/master/L16-gan/code/07_dcgan-cats-and-dogs.ipynb

Cats vs Dogs

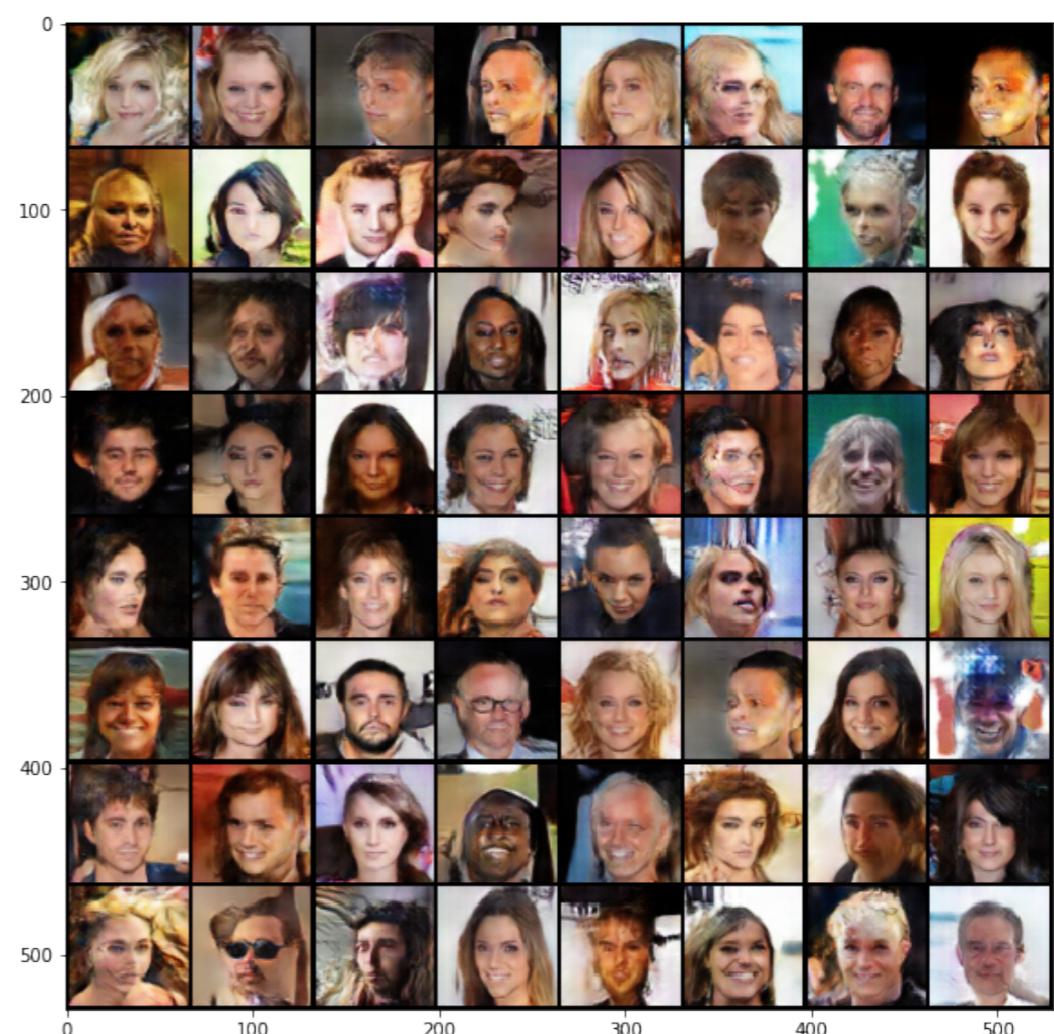
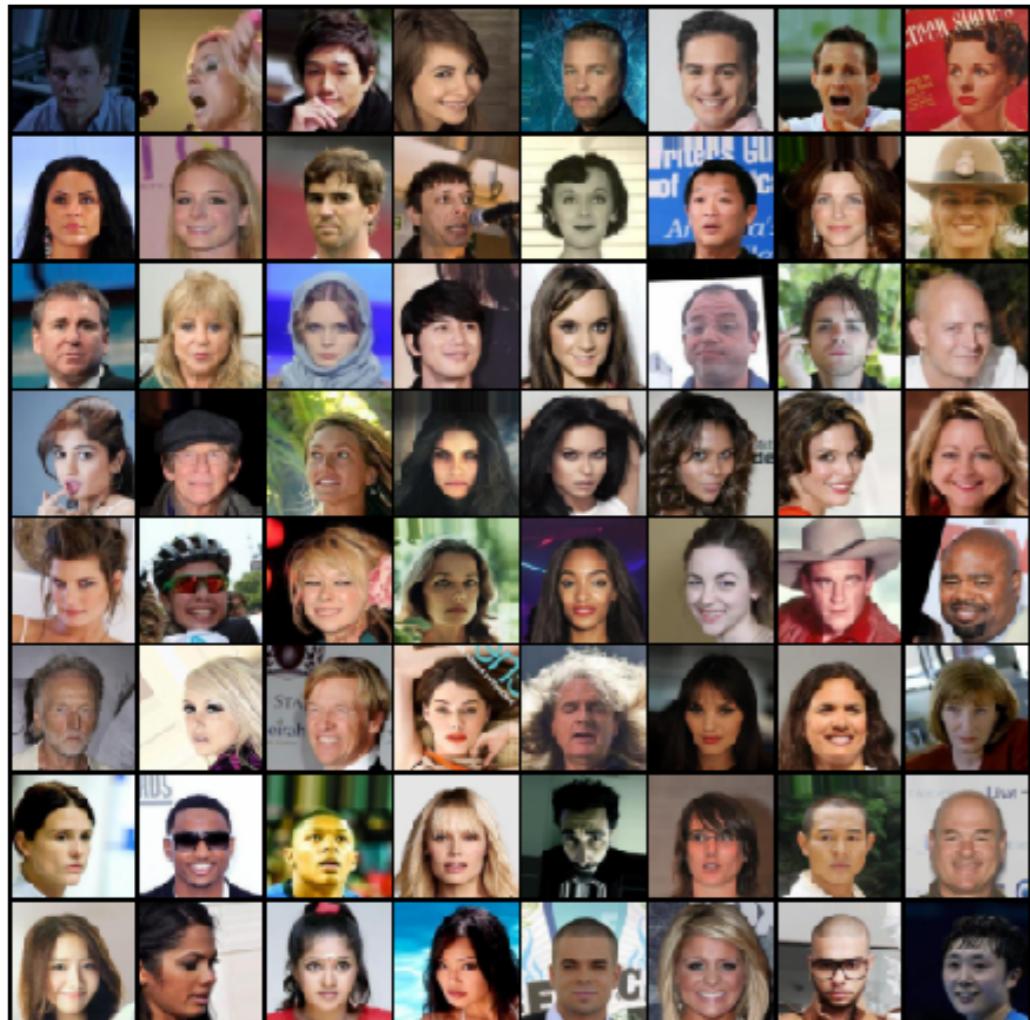
Training Images



https://github.com/rasbt/stat453-deep-learning-ss20/blob/master/L16-gan/code/07_dcgan-cats-and-dogs.ipynb

CelebA Face Images

Training Images



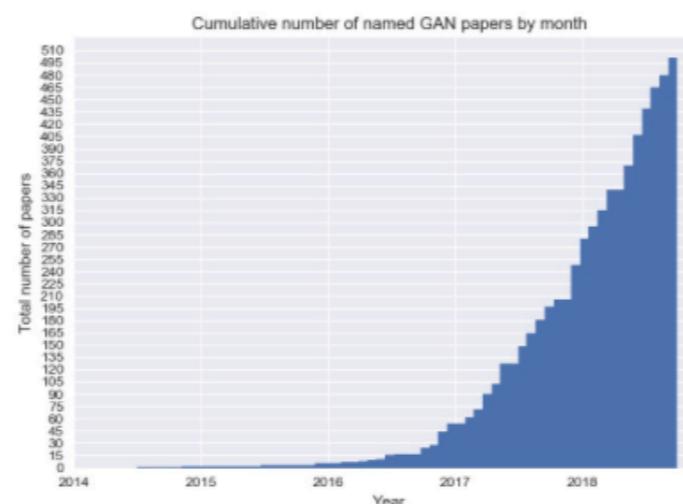
https://github.com/rasbt/stat453-deep-learning-ss20/blob/master/L16-gan/code/08_dcgan-celeba.ipynb

Lecture Overview

1. The Main Idea Behind GAN
2. The GAN Training
3. GANs in Practice
- 4. GAN Resources**

There are many more flavors of GANs that we have time to cover :)

A tentative list of named GANs <https://github.com/hindupuravinash/the-gan-zoo>



You can also check out the same data in a tabular format with functionality to filter by year or do a quick search by title [here](#).

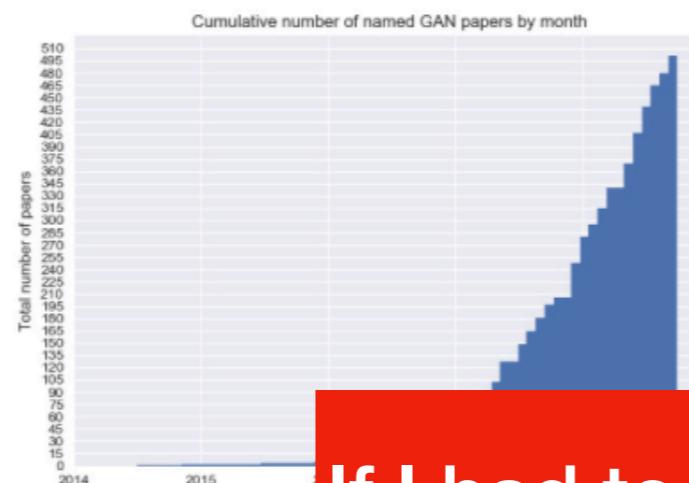
Contributions are welcome. Add links through pull requests in gans.tsv file in the same format or create an issue to lemme know something I missed or to start a discussion.

Check out [Deep Hunt](#) - my weekly AI newsletter for this repo as [blogpost](#) and follow me on [Twitter](#).

- 3D-ED-GAN - [Shape Inpainting using 3D Generative Adversarial Network and Recurrent Convolutional Networks](#)
- 3D-GAN - [Learning a Probabilistic Latent Space of Object Shapes via 3D Generative-Adversarial Modeling \(github\)](#)
- 3D-IWGAN - [Improved Adversarial Systems for 3D Object Generation and Reconstruction \(github\)](#)
- 3D-PhysNet - [3D-PhysNet: Learning the Intuitive Physics of Non-Rigid Object Deformations](#)
- 3D-RecGAN - [3D Object Reconstruction from a Single Depth View with Adversarial Learning \(github\)](#)
- ABC-GAN - [ABC-GAN: Adaptive Blur and Control for improved training stability of Generative Adversarial Networks \(github\)](#)
- ABC-GAN - [GANs for LIFE: Generative Adversarial Networks for Likelihood Free Inference](#)
- AC-GAN - [Conditional Image Synthesis With Auxiliary Classifier GANs](#)

There are many more flavors of GANs that we have time to cover :)

A tentative list of named GANs <https://github.com/hindupuravinash/the-gan-zoo>



You can also check out the same data in a tabular format w...

Contributions are welcome. Add links through pull requests if you know something I missed or to start a discussion.

Check out [Deep Hunt](#) - my weekly AI newsletter for this re...

- 3D-ED-GAN - [Shape Inpainting using 3D Generative Adversarial Networks](#)
- 3D-GAN - [Learning a Probabilistic Latent Space of Object Shapes via 3D Generative-Adversarial Modeling](#) ([github](#))
- 3D-IWGAN - [Improved Adversarial Systems for 3D Object Generation and Reconstruction](#) ([github](#))
- 3D-PhysNet - [3D-PhysNet: Learning the Intuitive Physics of Non-Rigid Object Deformations](#)
- 3D-RecGAN - [3D Object Reconstruction from a Single Depth View with Adversarial Learning](#) ([github](#))
- ABC-GAN - [ABC-GAN: Adaptive Blur and Control for improved training stability of Generative Adversarial Networks](#) ([github](#))
- ABC-GAN - [GANs for LIFE: Generative Adversarial Networks for Likelihood Free Inference](#)
- AC-GAN - [Conditional Image Synthesis With Auxiliary Classifier GANs](#)

If I had to pick 3, I would highlight

- Wasserstein GAN
- Cycle GAN
- CGAN (conditional GAN)

Further Reading Suggestions

- Original GAN paper
Goodfellow, Ian, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. "[Generative Adversarial Nets.](#)" In *Advances in Neural Information Processing Systems*, pp. 2672-2680. 2014.
- Salimans, Tim, Ian Goodfellow, Wojciech Zaremba, Vicki Cheung, Alec Radford, and Xi Chen. "[Improved techniques for training GANs.](#)" In *Advances in Neural Information Processing Systems*, pp. 2234-2242. 2016.
- Open Questions about Generative Adversarial Networks by Augustus Odena, <https://distill.pub/2019/gan-open-problems/>
- How to Train a GAN? Tips and Tricks to Make GANs Work by Soumith Chintala (PyTorch Dev.) <https://github.com/soumith/ganhacks>

Stuff in the News

SustaiNLP 2020

First Workshop on Simple and Efficient Natural Language Processing

<https://sites.google.com/view/sustainlp2020/shared-task>

SustaiNLP 2020 (co-located with EMNLP2020) is organizing a shared task to **promote the development of effective, energy-efficient models for difficult NLU tasks.**

As language models become crazily big these days, it's a good idea to base benchmarks not only on predictive performance but also factor in energy-efficiency -- hoping to see similar approaches for computer vision as well:

Learning Oracle Attention for High-fidelity Face Completion

<https://arxiv.org/abs/2003.13903>

Tong Zhou¹ Changxing Ding¹ Shaowen Lin¹ Xinchao Wang² Dacheng Tao³

¹ South China University of Technology ² Stevens Institute of Technology

³ UBTECH Sydney AI Centre, School of Computer Science, Faculty of Engineering,
The University of Sydney, Darlington, NSW 2008, Australia

201821011282@mail.scut.edu.cn chxding@scut.edu.cn

eeswlin@mail.scut.edu.cn xinchao.w@gmail.com dacheng.tao@sydney.edu.au

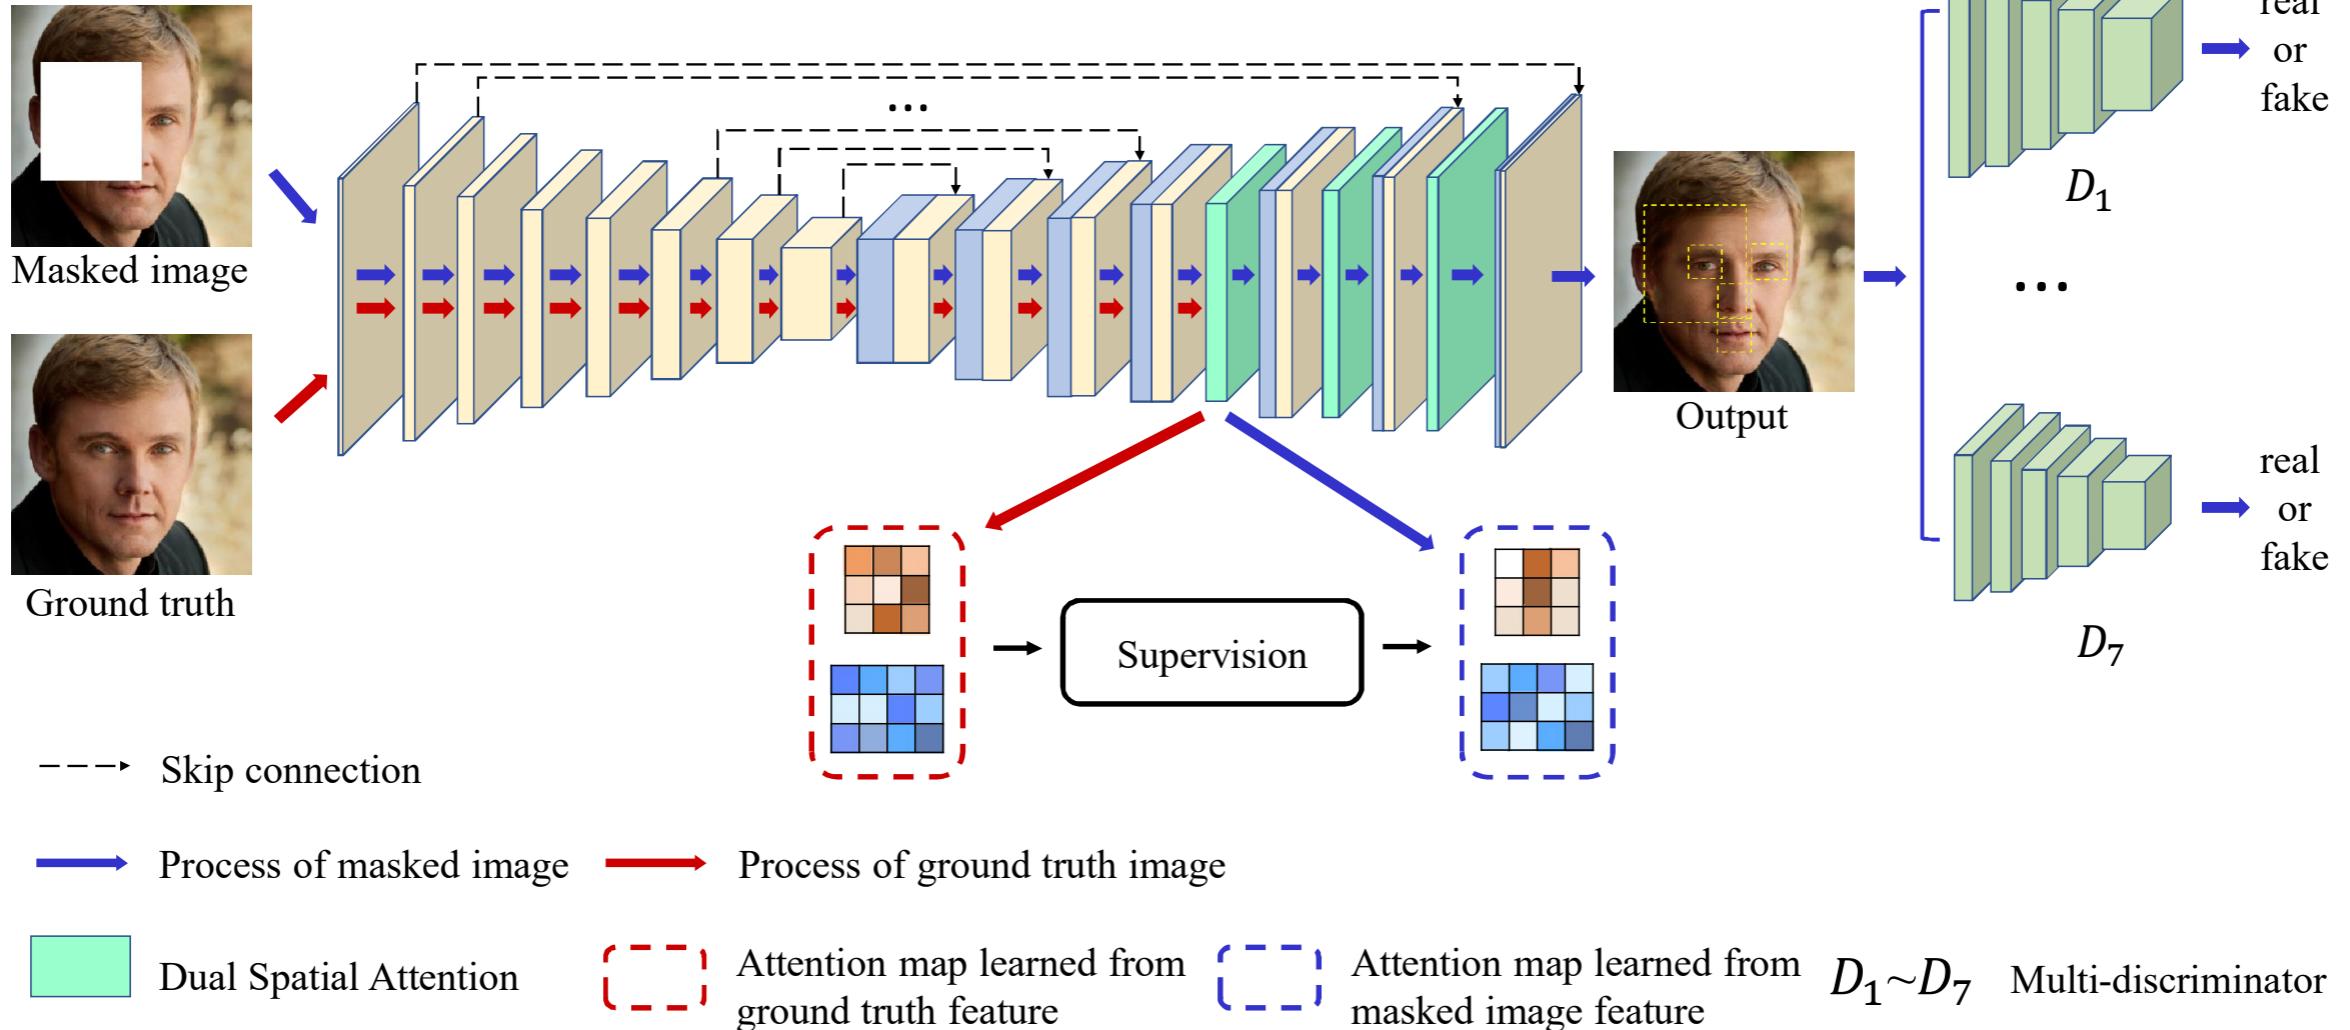


Figure 2. The overall architecture of our model. We use the U-Net structure [14] as backbone. The proposed attention module is embedded to layer 12, 13 and 14, where the resolution is 16×16 , 32×32 and 64×64 respectively. We feed the ground truth images to the network, such that we can impose an oracle supervision signal on attention scores produced by DSA. We deploy discriminators on both the masked area and each of the four facial components, respectively. Best viewed in color.



AI against coronavirus spread

Submit a photo from a room, hall or any other place where people might be in.

An Artificial Intelligence will recommend surfaces which require frequent disinfection to avoid contamination.

Select an image



How it works?

We use a semantic segmentation Deep Convolutional Neural Network to highlight objects prone to cross-contamination such as tables, doors, glasses and jars. The list of objects considered right now comprises 50+ categories.